

Relax! The Semilenient Core of Choreographic Programming (Functional Pearl)

DAN PLYUKHIN, University of Southern Denmark, Denmark

XUEYING QIN, University of Southern Denmark, Denmark

FABRIZIO MONTESI, University of Southern Denmark, Denmark

The past few years have seen a surge of interest in choreographic programming, a programming paradigm for concurrent and distributed systems. The paradigm allows programmers to implement a distributed interaction protocol with a single high-level program, called a *choreography*, and then mechanically *project* it into correct implementations of its participating processes. A choreography can be expressed as a λ -term parameterized by constructors for creating data “at” a process and for communicating data between processes. Through this lens, recent work has shown how one can add choreographies to mainstream languages like Java, or even embed choreographies as a DSL in languages like Haskell and Rust. These new choreographic languages allow programmers to write in applicative style (like in functional programming) and write higher-order choreographies for better modularity. But the semantics of functional choreographic languages is not well-understood. Whereas typical λ -calculi can have their operational semantics defined with just a few rules, existing models for *choreographic* λ -calculi have *dozens* of complex rules and *no clear or agreed-upon evaluation strategy*.

We show that functional choreographic programming is simple. Beginning with the Chor λ model from previous work, we strip away inessential features to produce a “core” model called λ^X . We discover that underneath Chor λ ’s apparently ad-hoc semantics lies a close connection to non-strict λ -calculi; we call the resulting evaluation strategy *semilenient*. Then, inspired by previous non-strict calculi, we develop a notion of *choreographic evaluation contexts* and a special *commute* rule to simplify and explain the unusual semantics of functional choreographic languages. The extra structure leads us to a presentation of λ^X with just ten rules, and a discovery of three missing rules in previous presentations of Chor λ . We also show how the extra structure comes with nice properties, which we use to simplify the correspondence proof between choreographies and their projections. Our model serves as both a principled foundation for functional choreographic languages and a good entry point for newcomers.

CCS Concepts: • **Theory of computation** → **Lambda calculus**; **Distributed computing models**; • **Computing methodologies** → **Distributed programming languages**.

Additional Key Words and Phrases: Choreographies, Concurrency, λ -calculus

ACM Reference Format:

Dan Plyukhin, Xueying Qin, and Fabrizio Montesi. 2025. Relax! The Semilenient Core of Choreographic Programming (Functional Pearl). *Proc. ACM Program. Lang.* 9, ICFP, Article 269 (August 2025), 27 pages. <https://doi.org/10.1145/3747538>

1 Introduction

Choreographic programming [Montesi 2013] is a paradigm for writing concurrent code. Programmers can write a single program, called a *choreography*, and *project* it (i.e., compile it) to

Authors’ Contact Information: Dan Plyukhin, University of Southern Denmark, Denmark, dplyukhin@imada.sdu.dk; Xueying Qin, University of Southern Denmark, Denmark, xyqin@imada.sdu.dk; Fabrizio Montesi, University of Southern Denmark, Denmark, fmontesi@imada.sdu.dk.



This work is licensed under a Creative Commons Attribution 4.0 International License.

© 2025 Copyright held by the owner/author(s).

ACM 2475-1421/2025/8-ART269

<https://doi.org/10.1145/3747538>

generate correct implementations of each process in the application. An oft-repeated slogan is “deadlock-freedom by design”, which means that processes projected from a choreography do not deadlock [Carbone and Montesi 2013]. This property is no accident: in fact, proof normalisation in linear logic corresponds to a first-order choreographic programming language [Carbone et al. 2018]. Choreographies also have many other practical benefits: in just the past four years, choreographies have been used to reduce proof burden for Hoare-style verification of concurrent systems [Cruz-Filipe et al. 2023a; van den Bos and Jongmans 2023], as an intermediate representation for distributed cryptography [Acay et al. 2021], and as a tool for implementing distributed applications [Lugovic and Montesi 2024]. For a comprehensive overview, we refer the reader to Montesi’s textbook [Montesi 2023].

A recent flurry of activity can be traced to 2020, when the first choreographic programming language for realistic software development appeared: *Choral* [Giallorenzo et al. 2020, 2024] showed how a mainstream language like Java can be made “choreographic” by adding datatypes for located values $T@p$, and by adding functions of type $T@p \rightarrow T@q$ to denote communication from process p to process q . Extending this idea, it was found that (at the cost of some expressive power) choreographic programming can be implemented as a library in sufficiently expressive languages [Shen et al. 2023]. These discoveries led to an explosion of choreographic programming in languages like Haskell [Bates et al. 2025; Shen et al. 2023], Rust [Bates et al. 2025; Laddad et al. 2024], Clojure [klo 2025], and Elixir [cho 2025]. Choral in particular has been used to develop practical applications like IRC [Lugovic and Montesi 2024] and model-serving pipelines [Plyukhin et al. 2024] with similar performance to hand-written processes and deadlock-freedom by design.

Many new choreographic languages are *higher-order*, meaning programmers can write choreographies that take other choreographies as parameters. Higher-order choreographies are useful for writing modular systems: for example, one can implement a replicated key-value store as a choreography and parameterize it by another choreography that implements the replication protocol [Shen et al. 2023]. Higher-order choreographic programming was first presented in Choral, and subsequently came attempts to formalize its semantics [Cruz-Filipe et al. 2022; Cruz-Filipe et al. 2023; Graversen et al. 2024; Hirsch and Garg 2022]. These formalizations all build on the λ -calculus, but unfortunately they lack its elegant simplicity. For instance, a semantics for Plotkin’s call-by-value λ -calculus can be given with just one axiom [Plotkin 1975], and the call-by-need λ -calculus only needs three [Ariola et al. 1995]. In contrast, choreographic λ -calculi currently use more than 20, sometimes almost 40, rules [Cruz-Filipe et al. 2022; Cruz-Filipe et al. 2023; Hirsch and Garg 2022].

Part of the complexity in past models stems from ambition. For instance, $\text{Chor}\lambda$ [Cruz-Filipe et al. 2022] includes constructors for algebraic datatypes, so it is not truly a “core” model. But more fundamentally, *researchers have not settled on an evaluation strategy!* Consider the two current foundational models:

- Pirouette [Hirsch and Garg 2022] is call-by-value, but this requires global synchronization—even for non-involved participants—on every choreography call. This is inconsistent with Pirouette’s intraprocedural semantics, which has more concurrency than call-by-value. Thus Pirouette lacks a reasonable version of the (β) axiom, which states that any expression can be factored out into a separate definition (or, dually, inlined) without changing the program’s semantics [Barendregt 1984].
- $\text{Chor}\lambda$ [Cruz-Filipe et al. 2022] is a more faithful model of higher-order choreographic languages, but its evaluation strategy is mysterious—in fact we shall see it is neither strict nor lazy. The model also crucially depends on ad-hoc “restructuring” rules, and its correctness

proof is quite intimidating. More recent models avoid these restructuring rules by ignoring recursion [Bates et al. 2025] or relying on extra synchronization [Graversen et al. 2024].

Our goal is to show that Chor λ 's approach is the right way to go, and that its unusual semantics is not a wart—it is a beauty mark. We make our case by developing a new presentation of the core of Chor λ and showing that it has a clear evaluation strategy combining features from strict (call-by-value) and non-strict (lenient, or call-by-future [Arvind et al. 1986]) calculi. This connection to non-strict calculi is particularly surprising because our semantics “emerges from” the semantics of the network, which is call-by-value. Because of its close connection to lenient calculi, we call our evaluation strategy *semilenient*.

The remainder of the paper presents λ^X , a model that reveals the elegant functional core at the heart of choreographic programming. Our key contributions are:

- *A streamlined model.* The operational semantics of λ^X has just ten rules. We accomplish this by cutting away inessential features (like Chor λ 's data structures and Pirouette's multiple abstractions) and by introducing choreographic evaluation contexts to capture the semilenient evaluation strategy. These evaluation contexts are governed by simple laws, which researchers can use as a recipe to find the right semantics in their own models. We use the extra structure to define a simple *commute* rule, which summarizes what would otherwise be eight seemingly ad-hoc rules. Using the new rule, we discover *three missing rules* in Chor λ 's published semantics that are necessary for choreographies to match the behavior of their projections.
- *A simplified correspondence proof.* The hallmark result of choreographic programming languages is a *Projection Theorem*, which explains how choreographies and networks correspond. With prior approaches, this result could only be proved by a large and difficult argument by structural induction. Doing so does not give much intuition about how choreographies relate to their projection, and it is easy to make mistakes—leading to incorrect definitions. Here we find that evaluation contexts can shed some light: it turns out that evaluation contexts in the choreography are projected into evaluation contexts at the network. We use this result, along with some other informative lemmas, to give a nice visual proof of the Projection Theorem by commuting diagrams.
- *A stronger correspondence result.* Chor λ has conspicuous features at the network level that are not present in ordinary process languages. When compiling to a conventional call-by-value language that lacks these features, it is unclear if important results like deadlock-freedom will actually hold. We fill in the missing piece with a novel *prophecy relation* that lets networks predict the future, and a *Prophecy Theorem* that shows networks with prophecy are no more powerful than regular networks. By working modulo prophecy, our Projection Theorem establishes a weak bisimulation between choreographies and plain old call-by-value networks.

Our model shows higher-order choreographic programming has a simple and elegant foundation in the λ -calculus. We believe these developments will serve as a good introduction for researchers to the beauty of choreographic programming, and as a practical jumping-off point for future work. We will proceed in three easy pieces. Section 2 introduces the network language—our compilation target—and explains the need for choreographies. Section 3 presents λ^X and its connections to other λ -calculi. Section 4 defines the projection from λ^X to the network level and a new technique for proving its correctness. We wrap up with related work in Section 5 and conclusions in Section 6. In some cases, proofs have been omitted for space; they can be found in the Supplemental Material.

2 Networks

Figure 1 presents our network language. We assume an unbounded set of *variables* x, y, z, \dots , of *procedure names* f, g, h, \dots , of *labels* l_1, l_2, \dots , and *process names* p, q, r, \dots . A value L may be a

Terms:

$$\begin{aligned}
 P, Q, R &::= L \mid P P \mid \text{if } P \text{ then } P \text{ else } P \mid \oplus_p l P \mid \&_p \{l_1 : P_1, \dots, l_n : P_n\} \mid f(\bar{p}) \\
 L &::= x \mid \lambda x : S. P \mid \text{send}_p \mid \text{recv}_q \mid c \mid \perp \\
 S &::= S \rightarrow S \mid \alpha \mid \perp \\
 \mathbb{P} &::= \{f_i(\bar{p}) : S_i = P_i\}_{i \in I} \\
 \mathcal{N} &::= p[P] \mid (\mathcal{N} \mid \mathcal{N})
 \end{aligned}$$

Frames:

$$\mathcal{F} ::= \bullet \mid P \mid L \bullet \mid \text{if } \bullet \text{ then } P \text{ else } P$$

Evaluation contexts:

$$\mathcal{E} ::= \bullet \mid \mathcal{F}[\mathcal{E}]$$

Notions of reduction:

$$\begin{aligned}
 \text{send}_p c &\xrightarrow{\text{send}_p c} \perp && (\text{send}) \\
 \text{recv}_p \perp &\xrightarrow{\text{recv}_p c} c && (\text{receive}) \\
 \oplus_p l P &\xrightarrow{\oplus_p l} P && (\text{choice}) \\
 \&_p \{l_1 : P_1, \dots, l_n : P_n\} &\xrightarrow{\&_p l_i} P_i && (\text{offer}) \\
 (\lambda x : T. P) L &\xrightarrow{\quad} P[x := L] && (p\text{-app}) \\
 \text{if } c \text{ then } P_{\text{true}} \text{ else } P_{\text{false}} &\xrightarrow{\quad} P_c && (p\text{-if}) \\
 f(\bar{p}) &\xrightarrow{\quad} P[\bar{q} := \bar{p}] && (p\text{-def}) \\
 &\quad \text{where } (f(\bar{q}) : S = P) \in \mathbb{P} \\
 \perp \perp &\xrightarrow{\quad} \perp && (\text{bottom})
 \end{aligned}$$

Evaluation strategy:

$$\begin{aligned}
 &\frac{P \mapsto P'}{p[\mathcal{E}[P]] \mid \mathcal{N} \xrightarrow{\tau} p[\mathcal{E}[P']] \mid \mathcal{N}} [p\text{-tau}] \\
 &\frac{P_1 \xrightarrow{\text{send}_q c} P'_1 \quad P_2 \xrightarrow{\text{recv}_p c} P'_2}{p[\mathcal{E}_1[P_1]] \mid q[\mathcal{E}_2[P_2]] \mid \mathcal{N} \xrightarrow{\text{comp}_{p,q}} p[\mathcal{E}_1[P'_1]] \mid q[\mathcal{E}_2[P'_2]] \mid \mathcal{N}} [p\text{-com}] \\
 &\frac{P_1 \xrightarrow{\oplus_p l} P'_1 \quad P_2 \xrightarrow{\&_q l_i} P'_2}{p[\mathcal{E}_1[P_1]] \mid q[\mathcal{E}_2[P_2]] \mid \mathcal{N} \xrightarrow{\text{select}_{p,q} l} p[\mathcal{E}_1[P'_1]] \mid q[\mathcal{E}_2[P'_2]] \mid \mathcal{N}} [p\text{-select}]
 \end{aligned}$$

Fig. 1. Syntax and semantics for networks

variable x , an abstraction $\lambda x : S. P$, a communication primitive send_p or recv_p , or a constant c . We assume the language includes constants **true** and **false** with the usual meaning, and a distinguished constant \perp that indicates “nothing left to do”. Note that process names and procedure names are not values; our choreography model will require a strict separation between processes and ordinary

$$\begin{array}{c}
\frac{p \in \Theta \quad \Theta; \Gamma \vdash P : S}{\Theta; \Gamma \vdash \oplus_p l P : S} [\text{NTCHO}] \quad \frac{p \in \Theta \quad \Theta; \Gamma \vdash P_i : S \text{ for } 1 \leq i \leq n}{\Theta; \Gamma \vdash \&_p \{l_1 : P_1, \dots, l_n : P_n\} : S} [\text{NTOFF}] \\
\\
\frac{p \in \Theta}{\Theta; \Gamma \vdash \text{send}_p : S \rightarrow \perp} [\text{NTSEND}] \quad \frac{p \in \Theta}{\Theta; \Gamma \vdash \text{recv}_p : \perp \rightarrow S} [\text{NTRCV}] \\
\\
\frac{\Gamma, x : S \vdash P : S'}{\Theta; \Gamma \vdash \lambda x : S. P : S \rightarrow S'} [\text{NTABS}] \quad \frac{x : S \in \Gamma}{\Theta; \Gamma \vdash x : S} [\text{NTVAR}] \\
\\
\frac{\Theta; \Gamma \vdash P_1 : S \rightarrow S' \quad \Theta; \Gamma \vdash P_2 : S}{\Theta; \Gamma \vdash P_1 P_2 : S'} [\text{NTAPP}] \\
\\
\frac{\Theta; \Gamma \vdash P : \text{Bool} \quad \Theta; \Gamma \vdash P_1 : S \quad \Theta; \Gamma \vdash P_2 : S}{\Theta; \Gamma \vdash \text{if } P \text{ then } P_1 \text{ else } P_2 : S} [\text{NTIF}] \\
\\
\frac{\text{type}(c) = \alpha}{\Theta; \Gamma \vdash c : \alpha} [\text{NTCONST}] \quad \frac{(f(\bar{q}) : S = P) \in \mathbb{P} \quad |\bar{p}| = |\bar{q}| \quad \text{distinct}(\bar{p}) \quad \bar{p} \subseteq \Theta}{\Theta; \Gamma \vdash f(\bar{p}) : S} [\text{NTDEF}]
\end{array}$$

Fig. 2. Typing rules for processes.

data, and combining the two requires process polymorphism [Graversen et al. 2024] which is beyond the scope of our work.

A *process* P is composed of familiar constructs like values L , applications $P_1 P_2$, and if-expressions *if* P' *then* P_1 *else* P_2 . We also include primitives for distributed choice, drawn from concurrency theory: $\&_p \{l_1 : P_1, \dots, l_n : P_n\}$ and $\oplus_p l P$, explained below. A *network* \mathcal{N} is a fixed set of processes $p_1[P_1] \mid \dots \mid p_n[P_n]$, where the process names p_1, \dots, p_n are distinct and the processes P_1, \dots, P_n are *closed*, i.e., they have no free variables.

To write non-terminating programs, the semantics is parameterized by a set of recursive procedure definitions $\mathbb{P} = \{f_i(\bar{p}) : S_i = P_i\}_{i \in I}$. Each procedure f_i in \mathbb{P} is parameterized by a list of process names \bar{p} . When $(f(\bar{p}) : S = P) \in \mathbb{P}$ and some process makes a procedure call $f(\bar{q})$, the process names \bar{q} are substituted for \bar{p} in the procedure body P . Thus procedures can be applied to process names and abstractions can be applied to values, but not vice versa.

For the static semantics of processes, we assume a set of base types $\alpha \in \{\text{Bool}, \perp, \dots\}$ and a function ‘type(–)’ mapping constants to their types. Typing judgments for processes have the form $\Theta; \Gamma \vdash_P P : S$, where Θ is a set of process names in scope, Γ is a typing context mapping from variables to types, \mathbb{P} is a set of procedure definitions, P is a process, and S is its type; for readability, we often leave \mathbb{P} implicit and instead write $\Theta; \Gamma \vdash P : S$. A set of procedure definitions \mathbb{P} is *well-typed* if, for each definition $f(\bar{p}) : S = P$ in \mathbb{P} , we have $\bar{p}; \Gamma \vdash P : S$. The full set of typing rules for processes appears in Figure 2.

A process can send q a value L by applying send_q to L ; likewise, it can wait for a message from p with the application $\text{recv}_p \perp$. A process can signal a control flow decision to q with the expression $\oplus_q l P$, which means “send label l to q and continue as P ”; conversely, it can wait for a decision from p with the expression $\&_p \{l_1 : Q_1, \dots, l_n : Q_n\}$, which means “upon receiving label l_i from p , continue as Q_i ”. Strictly speaking, distributed choice can be implemented with ordinary send_q and recv_p —but it is common practice to make the two primitives distinct in process calculi based on linear logic and session types. We include choice primitives to retain coherence with that work.

We give the network language a standard call-by-value semantics using evaluation contexts, with a small creative choice in how we present the latter. Usually, evaluation contexts are defined

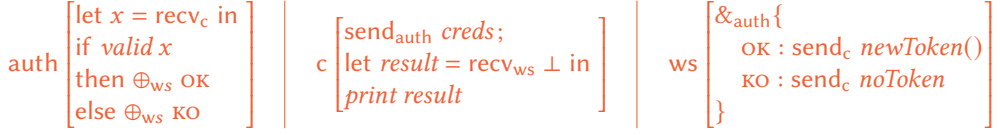


Fig. 3. A simple distributed authentication protocol based on OpenID [Montesi 2023]. The example uses let-sugar, i.e. $\text{let } x = P_1 \text{ in } P_2 \equiv (\lambda x. P_2) P_1$, and (;)-sugar, i.e. $P_1; P_2 \equiv (\lambda x. P_2) P_1$ where x is fresh. We will use this convention in the network language throughout the paper.

```

1   let x = comc,auth creds@c in
2   let result =
3       if valid@auth x then
4           selectauth,ws OK
5           newToken(ws)
6       else
7           selectauth,ws KO
8           noToken@ws
9   in print@c (comws,c result)

```

Fig. 4. A choreographic implementation of the distributed authentication protocol. Projecting this choreography generates a network equivalent to the one in Figure 3.

by direct induction like so:

$$\mathcal{E}' ::= \mathcal{E}' P \mid L \mathcal{E}' \mid \text{if } \mathcal{E}' \text{ then } P \text{ else } P \mid \bullet$$

Instead we have defined evaluation contexts as a stack of *frames* \mathcal{F} , so any evaluation context \mathcal{E} can be written as $\mathcal{F}_1[\mathcal{F}_2[\dots]]$. This presentation—which we borrowed from System F_J [Maurer et al. 2017], but dates back at least to Huet’s zippers [Huet 1997]—will become useful in Section 3 when we introduce choreographic evaluation contexts as stacks of both *choreographic frames* and so-called *answering contexts*. By routine induction, the definitions of \mathcal{E} and \mathcal{E}' are equivalent.

2.1 Motivating Choreographies

Figure 3 implements a simple authentication protocol in the network language, with types omitted for readability. Before reading our explanation below, we encourage the reader to stop and think about what will happen when Figure 3 is executed. Can the network reach a deadlock, i.e., a state where some processes are not values and yet the network has no next step?

Many programmers will intuitively reason about networks by sketching a sequence diagram [Object Management Group 2017] or using Alice-and-Bob notation [Needham and Schroeder 1978]. Others might construct a multiparty session type [Honda et al. 2016] and verify the network conforms to the type. But diagrams and session types both involve *supplementing* an existing network with extra information to specify its emergent behavior and *verifying* that the network actually respects the specification.

Choreographic programming is a more direct approach that aims to turn Alice-and-Bob notation into a *bona fide* programming language. This allows us to write concurrent programs at a higher level of abstraction and generate processes that implement the desired protocol. Figure 4 shows an

implementation of the authentication protocol as a λ^X choreography, with types omitted; we will not formalize its syntax and semantics until Section 3, so observe just the broad features for now. Notice the choreography looks like an ordinary functional program, but with three new attributes: (1) a syntax for tagging constants with locations, e.g. $\text{creds}@c$ represents the creds constant at process c ; (2) a primitive $\text{com}_{p,q}$ for communicating a value from p to q ; and (3) a primitive $\text{select}_{p,q} l$ so p can inform q about a control flow decision. Readers who drew a sequence diagram for Figure 3 can check that Figure 4 likely resembles their informal specification.

2.2 Authentication Explained

Reading the choreography in Figure 4, the protocol begins with an expression $\text{com}_{c,\text{auth}} \text{creds}@c$ that communicates the value creds from c to auth . This single expression corresponds to the terms $\text{send}_{\text{auth}} \text{creds}$ and $\text{recv}_c \perp$ in Figure 3. Next, auth checks if the credentials are valid and sends a signal to ws —either $\text{select}_{\text{auth},\text{ws}} \text{OK}$ or $\text{select}_{\text{auth},\text{ws}} \text{KO}$. These expressions correspond to the choice auth makes—either $\oplus_{\text{ws}} \text{OK}$ or $\oplus_{\text{ws}} \text{KO}$ —and the offer $\&_{\text{auth}} \{ \dots \}$ made by ws in Figure 3. Finally, the result is sent to c and printed. Figure 4 can be projected to produce a network very much like Figure 3, with minor differences that will be evident in Section 4. By construction, the projected network is both type-safe and deadlock-free [Cruz-Filipe et al. 2023].

Readers familiar with choreographic programming will notice Figure 4 is written in applicative style, like a conventional functional language. The if-expression on lines 3–8 returns a value that we can bind to the variable result , and line 9 applies $\text{print}@c$ directly to the expression $\text{com}_{\text{ws},c} \text{result}$ without giving a name to the intermediate result. Applicative features were not available to programmers until the Choral [Giallorenzo et al. 2024] language was released, and most formal models still only allow a first-order imperative programming style [Montesi 2023; Plyukhin et al. 2024]. In the next section, we explore why functional choreographic programming is so elusive.

3 Choreographies

Let us now introduce our choreography model λ^X properly. A (choreographic) value V may be a variable x , an abstraction $\lambda x : T. M$, a communication primitive $\text{com}_{p,q}$, or a constant $c@p$ located at process p . A choreography M may be an application $M_1 M_2$, an if-expression $\text{if } M' \text{ then } M_1 \text{ else } M_2$, a selection $\text{select}_{p,q} l M'$, a choreographic procedure call $f(\bar{p})$, or a value V . Choreographies also have let-expressions $\text{let } x = M \text{ in } M'$, which (unlike in the network language) are not just sugar—we shall see why this is useful below. Aside from let-expressions, the syntax of λ^X is essentially the same as our network model, except the communication primitives $\text{send}_q, \text{recv}_p$ have been unified by $\text{com}_{p,q}$ and the choice primitives $\oplus_q, \&_p$ have been unified by $\text{select}_{p,q}$. We also assume if-expressions and choreographic procedure calls are either let-bound or in tail position; for example, $f(\bar{p}) M$ must be expanded as $\text{let } x = f(\bar{p}) \text{ in } M$. The syntax is summarized in Figure 5.

Choreographic types T are also similar to network-level types. Base types $\alpha@p$ are now tagged with their location, and function types $T \rightarrow_{\bar{p}} T'$ are tagged with a list of processes \bar{p} called *proxy processes*; our type system will ensure that if a process p is involved in computing the function's result, then p will occur in T, T' , or \bar{p} . If the function type has no proxy processes, we will write it as $T \rightarrow T'$. The set of processes that *occur* in a term is formalized by the function $\text{pn}(-)$, defined at the bottom of Figure 5.

Typing judgments for choreographies, written $\Theta; \Gamma \vdash_D M : T$, are analogous to typing judgments for processes: Θ is a set of process names in scope, Γ is a mapping from variables to choreographic types, D is a set of choreographic procedure definitions, M is a choreography, and T is a choreographic type; for readability we usually leave D implicit. The λ^X typing rules are defined in Figure 6 and are very similar to the network-level typing rules. Notice how the network-level rules [NTSEND], [NTRCV] have been unified by the choreographic rule [TCom]—likewise for

Terms:

$$\begin{aligned}
 M &::= V \mid M M \mid \text{let } x : T = M \text{ in } M \mid \text{if } M \text{ then } M \text{ else } M \mid \text{select}_{p,q} l M \mid f(\bar{p}) \\
 V &::= x \mid \lambda x : T. M \mid \text{com}_{p,q} \mid c@p \\
 T &::= T \rightarrow_{\bar{p}} T \mid \alpha@p \\
 \mathbb{D} &::= \{f_i(\bar{p}) : T_i = M_i\}_{i \in I}
 \end{aligned}$$

Frames:

$$\mathcal{F} ::= \bullet M \mid V \bullet \mid \text{if } \bullet \text{ then } M \text{ else } M \mid \text{let } x : T = \bullet \text{ in } M$$

Answering contexts:

$$\mathcal{A} ::= \text{let } x : T = M \text{ in } \bullet \mid \text{select}_{q,r} l \bullet$$

Evaluation contexts:

$$\mathcal{E}_{\bar{p}} ::= \bullet \mid \mathcal{F}[\mathcal{E}_{\bar{p}}] \mid \mathcal{A}[\mathcal{E}_{\bar{p}}] \quad \text{where } \bar{p} \# \text{pn}(\mathcal{A})$$

Notions of reduction:

$$\begin{aligned}
 \text{com}_{p,q} c@p &\xrightarrow{\text{com}_{p,q}} c@q && (com) \\
 \text{select}_{p,q} l M &\xrightarrow{\text{select}_{p,q} l} M && (select) \\
 (\lambda x : T. M) M' &\xrightarrow{\bar{p}} \text{let } x : T = M' \text{ in } M && (app) \\
 \text{let } x : T = V \text{ in } M &\xrightarrow{\bar{p}} M[x := V] && (let) \\
 &\quad \text{where } \bar{p} = \text{pn}(V) \\
 \text{if } c@p \text{ then } M_{\text{true}} \text{ else } M_{\text{false}} &\xrightarrow{p} M_c && (if) \\
 f(\bar{p}) &\xrightarrow{\bar{p}} M[\bar{q} := \bar{p}] && (def) \\
 &\quad \text{where } (f(\bar{q}) : T = M) \in \mathbb{D} \\
 \mathcal{F}[\mathcal{A}[M]] &\xrightarrow{\bar{p}} \mathcal{A}[\mathcal{F}[M]] && (commute)
 \end{aligned}$$

Evaluation strategy:

$$\begin{aligned}
 \frac{M \xrightarrow{\bar{p}} M' \quad p \in \bar{p}}{\mathcal{E}_p[M] \xrightarrow{\tau} \mathcal{E}_p[M']} [c\text{-tau}] & \quad \frac{M \xrightarrow{\text{com}_{p,q}} M'}{\mathcal{E}_{p,q}[M] \xrightarrow{\text{com}_{p,q}} \mathcal{E}_{p,q}[M']} [c\text{-com}] & \quad \frac{M \xrightarrow{\text{select}_{p,q} l} M'}{\mathcal{E}_{p,q}[M] \xrightarrow{\text{select}_{p,q} l} \mathcal{E}_{p,q}[M']} [c\text{-select}]
 \end{aligned}$$

Mentioned processes:

$$\begin{aligned}
 \text{pn}(M_1 M_2) &= \text{pn}(M_1) \cup \text{pn}(M_2) & \text{pn}(\text{select}_{p,q} l M) &= \{p, q\} \cup \text{pn}(M) \\
 \text{pn}(\text{if } M \text{ then } M_1 \text{ else } M_2) &= & \text{pn}(\text{let } x : T = M \text{ in } M') &= \\
 \text{pn}(M) \cup \text{pn}(M_1) \cup \text{pn}(M_2) & & \text{pn}(T) \cup \text{pn}(M) \cup \text{pn}(M') & \\
 \text{pn}(\lambda x : T. M) &= \text{pn}(T) \cup \text{pn}(M) & \text{pn}(x) &= \text{pn}(\text{type}(x)) \\
 \text{pn}(\text{com}_{p,q}) &= \{p, q\} & \text{pn}(c@p) &= \text{pn}(T@p) = p \\
 \text{pn}(\bullet) &= \emptyset & \text{pn}(T_1 \rightarrow_{\bar{p}} T_2) &= \text{pn}(T_1) \cup \text{pn}(T_2) \cup \bar{p}
 \end{aligned}$$

Fig. 5. Syntax and semantics for $\lambda\lambda^X$

$$\begin{array}{c}
\frac{\Theta'; \Gamma, x : T \vdash M : T' \quad \text{pn}(T \rightarrow_{\bar{p}} T') = \Theta' \subseteq \Theta}{\Theta; \Gamma \vdash \lambda x : T. M : T \rightarrow_{\bar{p}} T'} [\text{TAbs}] \\
\\
\frac{x : T \in \Gamma \quad \text{pn}(T) \subseteq \Theta}{\Theta; \Gamma \vdash x : T} [\text{TVar}] \quad \frac{\Theta; \Gamma \vdash M_1 : T \rightarrow_{\bar{p}} T' \quad \Theta; \Gamma \vdash M_2 : T}{\Theta; \Gamma \vdash M_1 M_2 : T'} [\text{TApp}] \\
\\
\frac{\Gamma \vdash M : \text{Bool}@p \quad \Theta; \Gamma \vdash M_1 : T \quad \Theta; \Gamma \vdash M_2 : T \quad p \in \Theta}{\Theta; \Gamma \vdash \text{if } M \text{ then } M_1 \text{ else } M_2 : T} [\text{TIf}] \\
\\
\frac{\Theta; \Gamma \vdash M : T \quad p, q \in \Theta}{\Theta; \Gamma \vdash \text{select}_{p,q} l M : T} [\text{TSel}] \\
\\
\frac{p \in \Theta \quad \text{type}(c) = \alpha}{\Theta; \Gamma \vdash c@p : \alpha@p} [\text{TConst}] \quad \frac{p, q \in \Theta}{\Theta; \Gamma \vdash \text{com}_{p,q} : \alpha@p \rightarrow_{\emptyset} \alpha@q} [\text{TCom}] \\
\\
\frac{(f(\bar{q}) : T = M) \in \mathbb{D} \quad |\bar{p}| = |\bar{q}| \quad \text{distinct}(\bar{p})}{\Theta; \Gamma \vdash f(\bar{p}) : T[\bar{q} := \bar{p}]} [\text{TDef}] \\
\\
\frac{\Theta; \Gamma \vdash M_1 : T_1 \quad \Theta; \Gamma, x : T_1 \vdash M_2 : T_2}{\Theta; \Gamma \vdash \text{let } x : T_1 = M_1 \text{ in } M_2 : T_2} [\text{TLet}]
\end{array}$$

Fig. 6. Full set of typing rules for λ^χ .

[NTCHO], [NTOFF] and [TSEL]. A set of choreographic procedure definitions \mathbb{D} is *well-typed* if, for each definition $f(\bar{p}) : T = M$ in \mathbb{D} , we have $\bar{p}; \Gamma \vdash_{\mathbb{D}} M : T$.

The key novelty of λ^χ is how we present its dynamic semantics. But before diving into the details, let us discuss evaluation strategies.

3.1 Evaluation Strategies

The dynamic semantics of a λ -calculus can be presented in two parts. The first part is a *notion of reduction* $M \xrightarrow{\mu} M'$, which defines atomic computational steps such as β -reduction; we say M is a redex, M' is its contractum, and μ is an optional transition label. The second part is an *evaluation strategy* $M \xrightarrow{\mu} M'$, which defines the order redexes should be reduced [Barendregt 1984]. Two well-known evaluation strategies are call-by-value (“strict”, like our network language) and call-by-need (“lazy”, elegantly presented in the call-by-need calculus [Ariola et al. 1995]). These two evaluation strategies are deterministic: for any process P in our network language there is at most one P' such that $P \xrightarrow{\mu} P'$. (However, a network N may have more than one N' such that $N \xrightarrow{\mu} N'$, because processes execute concurrently.)

Choreographic programming languages have nondeterministic evaluation strategies. To motivate this, consider the first-order choreography in Figure 7. It defines a choreographic procedure $\text{loop}(p, q)$ which passes the integer 1 between p and q *ad infinitum*. Processes r_1 and r_2 enter this procedure on line 4, and r_3 sends the integer 2 to r_4 on line 5. A typical projection for this choreography is shown on the right; notice that the loop procedure was projected into two implementations, one for each of its role parameters.

We expect choreographic programs to have the same behavior as their projection. Inspecting the projected network, we see that the execution of r_1, r_2 and r_3, r_4 can be arbitrarily interleaved. This implies that a usual sequential semantics is inadequate for the choreography in Figure 7: it must be possible for the procedure call on line 4 to execute *out of order* with the communication on line 5, because the instructions involve different processes. This semantics, where the semicolon

1	$loop(p, q) =_{def}$	$loop_p(q) =_{def}$	$loop_q(p) =_{def}$
2	$com_{p,q} 1@p;$	$send_q 1;$	$recv_p \perp;$
3	$loop(q, p)$	$loop_q(q)$	$loop_p(p)$
4	$loop(r_1, r_2);$		
5	$com_{r_3,r_4} 2$	$r_1[loop_p(r_2)] \mid r_2[loop_q(r_1)] \mid r_3[send_{r_4} 1] \mid r_4[recv_{r_3} \perp]$	

Fig. 7. A choreography with a loop (left) and a typical projection (right).

operator is not quite sequential and not quite parallel, is the *de facto* standard in modern choreography languages and textbooks [Montesi 2023] and is crucial for the linear logic interpretation of choreographies [Carbone et al. 2018].

3.2 Lenient Semantics

When adapting a first-order language to the higher-order setting, we expect certain equivalences to be obeyed. Functional languages often use let-sugar and (\cdot)-sugar, as seen in Figure 3. These desugarings give us a hint about what the meaning of an application $(\lambda x : T. M) M'$ should be. Specifically, we expect these three choreographies to have identical semantics:

$$(\lambda x. com_{r_3,r_4} 2) loop(r_1, r_2) \quad (1)$$

$$let x = loop(r_1, r_2) in com_{r_3,r_4} 2. \quad (2)$$

$$loop(r_1, r_2); com_{r_3,r_4} 2. \quad (3)$$

Seen in this light, one might suggest a *lenient* semantics. Lenient evaluation is a nondeterministic evaluation strategy that is neither strict nor lazy [Tremblay 2000] but has much in common with call-by-need [Ariola et al. 1995; Arvind et al. 1996]. Its famous exponents include Id [Arvind et al. 1986], parallel Haskell [Arvind et al. 1996], and the Verse calculus [Augustsson et al. 2023].

Lenient semantics allows expressions to be evaluated concurrently *up to data dependency*. Hence it would indeed allow $com_{r_3,r_4} 2$ and $loop(r_1, r_2)$ to execute concurrently because neither expression depends on the other. But unfortunately, a lenient semantics would be *too concurrent!* Specifically, it allows us to unfold the $loop(r_1, r_2)$ call once:

$$com_{r_1,r_2} 1@r_1; loop(r_2, r_1); com_{r_3,r_4} 2 \quad (4)$$

And then to unfold the $loop(r_2, r_1)$ call:

$$com_{r_1,r_2} 1@r_1; com_{r_2,r_1} 1@r_2; loop(r_1, r_2); com_{r_3,r_4} 2 \quad (5)$$

And then to reduce $com_{r_2,r_1} 1@r_2$. In other words, a lenient semantics would tell us r_2 can send a message to r_1 in the initial state of the network. But by inspecting the projected code in Figure 7, we can see that r_2 will never send a message until it first receives a message from r_1 .

Thus lenient evaluation captures the *concurrency up to data dependency* of choreographies, but fails to capture *process dependencies* that enforce sequentiality. Armed with this insight, we will use prior art in non-strict λ -calculi [Ariola et al. 1995; Arvind et al. 1996; Maurer et al. 2017] to develop a language that is lenient up to process ordering—or *semilenient* for short.

3.3 Semilenient Semantics

We now explain λ^X 's semilenient semantics formally. We do this by defining notions of reduction—rules (*com*), (*select*), (*app*), (*let*), (*if*), (*def*), and (*commute*) in Figure 5—and an appropriate notion of evaluation context. Unlike ordinary evaluation contexts, a choreographic evaluation context is

indexed by a set of process names $\mathcal{E}_{\bar{p}}$. Decomposing a choreography M with an evaluation context $M = \mathcal{E}_{\bar{p}}[M']$ means that, if $p \in \bar{p}$, then p will evaluate M' in its next step. Since choreographies are concurrent, a term M can have distinct decompositions $M = \mathcal{E}_p[M'] = \mathcal{E}_q[M'']$ for different processes p, q .

Our evaluation contexts may appear complex, but in fact there are principles we can use to derive them. Whereas evaluation contexts in the network language are stacks of frames (Section 2), choreographic evaluation contexts are stacks of (choreographic) frames together with so-called *answering contexts*, explained below.

3.3.1 Choreographic Frames and Answering Contexts. A (choreographic) frame \mathcal{F} is the choreographic analogue of a network-level frame. Notice that each kind of network-level frame in Figure 1 corresponds directly to a choreographic frame in Figure 5; the only extra case we add is $\mathcal{F} ::= \text{let } x : T = \bullet \text{ in } M$ because let-bindings are no longer just syntactic sugar in λ^X .

An *answering context* \mathcal{A} is a context where, if $p \notin \text{pn}(\mathcal{A})$, then p has no work left to do. By filling an answering context with a value $\mathcal{A}[V]$, we make a choreographic term where p is ready to pass V to its enclosing context. Answering contexts are analogous to the notion of *answers* in the call-by-need lambda calculus [Ariola et al. 1995] and the notion of *tail contexts* in System F_J [Maurer et al. 2017].

How did we arrive at these definitions for \mathcal{F} and \mathcal{A} ? In fact, they emerge as properties of the projection function, which we define in Section 4. Choreographic frames are defined so that \mathcal{F} projects to a network-level frame \mathcal{F} . Meanwhile, answering contexts are defined so that the projection of $\mathcal{A}[M]$ for p is equal to the projection of M whenever $p \notin \text{pn}(\mathcal{A})$. In other words, answering contexts are precisely the contexts that disappear during projection.

3.3.2 Notions of Reduction and the Commute Rule. Most of the notions of reduction in λ^X should be unsurprising. The (*com*) rule moves a constant from p to q . The (*select*) rule simply steps into its continuation; it is a synchronization between p and q . The (*if*) and (*def*) rules are the same as in the network level, but with extra annotations to involve processes. However, the network's (*p-app*) rule is split into two steps: (*app*) reduces an application $(\lambda x : T.M) M'$ into an explicit let-binding $\text{let } x : T = M' \text{ in } M$, allowing our evaluation contexts to begin evaluating M without waiting for M' to be a value. Once M' is a value, the (*let*) rule substitutes that value into the body of the let-binding. (Another example of this can be seen in the call-by-need calculus, where the context $\text{let } x = M' \text{ in } M$ denotes a “thunk” x bound to M' in the expression M [Ariola et al. 1995].) The transition labels for all these rules are generally determined by the redex, except in the case of (*app*) and (*commute*) where the label \bar{p} can be any list of process names.

The (*commute*) rule is unusual: it pushes a frame \mathcal{F} into an answering context \mathcal{A} . This rule, which we borrowed from System F_J [Maurer et al. 2017], precisely captures the mysterious and ad-hoc “restructuring” rules from Chor λ . Consider the term:

$$\left(\text{let } x = f(r) \text{ in } \right)_{\text{com}_{p,q}} 0@p,$$

where $f(r)$ diverges. Its projection at p is simply $\text{send}_q 0$ because the call $f(r)$ is irrelevant to p . (Our type system guarantees, in general, that $f(\bar{r})$ only involves \bar{r} .) Likewise, the projection at q is $\text{recv}_p \perp$. In fact we can construct many examples like this:

$$\text{com}_{p,q} \left(\text{let } x = f(r) \text{ in } \right)_{0@p} \quad \text{let } x = \left(\text{let } y = f(r) \text{ in } \right)_{0@p} \text{ in } \quad \text{com}_{p,q} \left(\text{if } \left(\text{let } x = f(r) \text{ in } \right)_{\text{true}@p} \text{ then } 0@p \text{ else } 1@p \right)$$

In each case above, the choreography semantics should allow p and q to communicate without waiting for r —but the nonterminating computation $f(r)$ prevents us from ever creating the redex $\text{com}_{p,q} 0@p$. Higher-order choreographic languages therefore need rules for rewriting terms to expose reducible expressions; let us call these extra rules *restructuring rules*.

How does one know what restructuring rules a choreography language will need? In past approaches, there was only one way to find out: try to prove that choreographies and networks correspond, and add more restructuring rules when the proof gets stuck. This solution is unappealing because even a simple language requires many restructuring rules, and it is easy to miss cases when working by hand. Our presentation reveals a hidden structure behind the restructuring rules: they are all instances of the (*commute*) rule.

Figure 8 shows the restructuring rules generated by (*commute*) if we inline the definitions of \mathcal{F} and \mathcal{A} . Some of these rules further reinforce the connection between choreographies and non-strict calculi: (*let-let*) and (*app-let*) correspond to (*let-A*) and (*let-C*) from the call-by-need λ -calculus [Ariola et al. 1995]. These rules, along with (*let-app*), (*app-sel*), and (*sel-app*), are all found (with a slightly different presentation) in Chor λ [Cruz-Filipe et al. 2023].

Surprisingly, the remaining three rules (*if-let*), (*if-sel*), and (*let-sel*) from Figure 8 are all missing from the original publication of Chor λ . Upon closer inspection, we discovered that all three rules are indeed necessary for the correspondence result, and they were missed because they correspond to subtle edge cases in the proof. Chor λ 's semantics tells us the following three programs will never evaluate M_1 if $f(r)$ diverges, but any compiler that uses Chor λ 's projection function *will* in fact evaluate M_1 :

$\text{if } \left(\begin{array}{l} \text{let } x = f(r) \text{ in} \\ \text{true}@p \end{array} \right)$	$\text{let } x = f(r) \text{ in}$	$\text{let } x = f(r) \text{ in}$
$\text{then } M_1 \text{ else } M_2$	$\text{if } (\text{select}_{r,s} l \text{ true}@p)$	$\text{let } y = (\text{select}_{r,s} l \text{ true}@p) \text{ in}$
	$\text{then } M_1 \text{ else } M_2$	$\text{if } y \text{ then } M_1 \text{ else } M_2$

The fact that the (*commute*) rule guided us to these missing rules in Chor λ is a testament to the usefulness of our approach.

3.4 Properties of λ^X

Applying (*app*) and (*commute*) reductions is similar to putting a term in Administrative Normal Form (ANF) [Flanagan et al. 1993]. We define normal forms like so:

Definition 3.1. A choreography M is in *normal form* if it cannot be expressed in the form $\mathcal{E}_p[\Delta]$ for any p , where Δ is a redex for (*app*) or (*commute*).

Every closed choreography has a normal form. In the following, $M \xrightarrow{\tau} M'$ denotes a sequence of zero or more τ transitions $M \xrightarrow{\tau} \dots \xrightarrow{\tau} M'$:

LEMMA 3.2. For any closed choreography M , there exists \tilde{M} in normal form where $M \xrightarrow{\tau} \tilde{M}$ using only (*commute*) and (*app*) reductions.

PROOF. Omitted for space. Complete proofs can be found in the Supplemental Material. \square

As we saw in Section 3.3.2, putting a choreography in normal form can expose redexes that allow processes to make progress. Normal forms are also simplify proofs, as we will see in Section 4.

Next, we establish an important lemma about the “next steps” a choreography can take. We will use the superscript $(-)^*$ to denote a stack of contexts—for example, \mathcal{A}^* is either a hole \bullet or an answering context filled with a stack $\mathcal{A}[\mathcal{A}^*]$.

Definition 3.3. A *redex at p* is a choreographic redex Δ where either (1) $\Delta \xrightarrow{\mu} \Delta'$ for some μ, Δ' where $p \in \text{pn}(\mu)$, (2) $\Delta = \text{com}_{q,p} M$ for some q, M , or (3) $\Delta = \text{if } M \text{ then } M_1 \text{ else } M_2$ for some M, M_1, M_2 where $p \in (\text{pn}(M_1) \cup \text{pn}(M_2)) \setminus \text{pn}(M)$.

$\text{if } (\text{let } x = M_1 \text{ in } M_2) \text{ then } M_3 \text{ else } M_4$	\mapsto	$\text{let } x = M_1 \text{ in } (\text{if } M_2 \text{ then } M_3 \text{ else } M_4)$	(if-let)
$\text{let } y = (\text{let } x = M_1 \text{ in } M_2) \text{ in } M_3$	\mapsto	$\text{let } x = M_1 \text{ in } (\text{let } y = M_2 \text{ in } M_3)$	(let-let)
$V \text{ (let } x = M_1 \text{ in } M_2)$	\mapsto	$\text{let } x = M_1 \text{ in } V M_2$	(app-let)
$(\text{let } x = M_1 \text{ in } M_2) M_3$	\mapsto	$\text{let } x = M_1 \text{ in } M_2 M_3$	(let-app)
$V \text{ (select}_{p,q} l M)$	\mapsto	$\text{select}_{p,q} l (V M)$	(app-sel)
$(\text{select}_{p,q} l M_1) M_2$	\mapsto	$\text{select}_{p,q} l (M_1 M_2)$	(sel-app)
$\text{if } (\text{select}_{p,q} l M_1) \text{ then } M_2 \text{ else } M_3$	\mapsto	$\text{select}_{p,q} l (\text{if } M_1 \text{ then } M_2 \text{ else } M_3)$	(if-sel)
$\text{let } x = \text{select}_{p,q} l M_1 \text{ in } M_2$	\mapsto	$\text{select}_{p,q} l (\text{let } x = M_1 \text{ in } M_2)$	(let-sel)

Fig. 8. This figure shows all the rules entailed by (*commute*) if we wrote them out explicitly.

LEMMA 3.4. Let M be a closed choreography with $p \in \text{pn}(M)$. Either:

- (1) $M = \mathcal{E}_p[\Delta]$ for some evaluation context \mathcal{E}_p where Δ is a redex at p , or
- (2) $M = \mathcal{A}^*[V]$ for some stack of answering contexts \mathcal{A}^* where $p \notin \text{pn}(\mathcal{A}^*)$ and V is a value.

The above result generalizes the usual property in deterministic languages, where every term is either a value or has an available redex. The next lemma implies that if M is in normal form, the decomposition $M = \mathcal{E}_p[\Delta]$ is unique for each process p .

LEMMA 3.5. Let M be a closed choreography in normal form and let p, q be processes that are not necessarily distinct. Assume $M = \mathcal{E}_p[\Delta_1] = \mathcal{E}_q[\Delta_2]$ for some evaluation contexts $\mathcal{E}_p, \mathcal{E}_q$ where Δ_1 and Δ_2 are redexes at both p and q . Then $\mathcal{E}_p = \mathcal{E}_q$ and $\Delta_1 = \Delta_2$.

We conclude by showing the type system for λ^X is sound with respect to our semantics. This result, coupled with the Projection Theorem, will allow us to prove that projections are deadlock-free in Section 4.

THEOREM 3.6 (PROGRESS OF EVALUATION). Let M be a choreography. If there exist Θ, T such that $\Theta; \emptyset \vdash M : T$, then either M is a value V or there exists M' such that $M \xrightarrow{\mu} M'$.

THEOREM 3.7 (PRESERVATION OF EVALUATION). Let M be a choreography. If there exist Θ, Γ, T such that $\Theta; \Gamma \vdash M : T$, then $\Theta; \Gamma \vdash M' : T$ for any M' such that $M \xrightarrow{\mu} M'$.

4 Projection

Now the rubber meets the road: we present the *projection* function, which compiles λ^X choreographies into networks. We will also show that any execution of the choreography can be matched by an execution of its projection (*completeness*) and any execution of the projection can be matched by the choreography (*soundness*). Together, these two properties are called the Projection Theorem.

Our proof strategy will take a significantly different route from prior work [Cruz-Filipe et al. 2023; Hirsch and Garg 2022; Montesi 2023]. We begin with a series of useful lemmas relating choreographic and network-level evaluation contexts. Then, instead of proving a direct correspondence between choreographies and their projections, we do it in two steps: first establishing a correspondence between choreographies and “superpowered” networks that can predict the future, and then showing that superpowered networks are equivalent to ordinary networks. The payoff for this extra work will be an elegant proof of the Projection Theorem and a clean correspondence (namely, a weak bisimulation) between λ^X and the network language.

$$\begin{aligned}
\llbracket \lambda x : T. M \rrbracket_p &= \begin{cases} \lambda x : \llbracket T \rrbracket_p. \llbracket M \rrbracket_p & \text{if } p \in \text{pn}(T) \cup \text{pn}(M) \\ \perp & \text{otherwise} \end{cases} \\
\llbracket \text{let } x : T = M_1 \text{ in } M_2 \rrbracket_p &= \begin{cases} (\lambda x : \llbracket T \rrbracket_p. \llbracket M_2 \rrbracket_p) \llbracket M_1 \rrbracket_p & \text{if } p \in \text{pn}(T) \cup \text{pn}(M_1) \\ \llbracket M_2 \rrbracket_p & \text{otherwise} \end{cases} \\
\llbracket M_1 M_2 \rrbracket_p &= \begin{cases} \llbracket M_1 \rrbracket_p \llbracket M_2 \rrbracket_p & \text{if } p \in \text{pn}(M_1) \cup \text{pn}(M_2) \\ \perp & \text{otherwise} \end{cases} \\
\llbracket \text{if } M \text{ then } M_1 \text{ else } M_2 \rrbracket_p &= \begin{cases} \text{if } \llbracket M \rrbracket_p \text{ then } \llbracket M_1 \rrbracket_p \text{ else } \llbracket M_2 \rrbracket_p & \text{if } p \in \text{pn}(\text{type}(M)) \\ (\lambda x : \perp. \llbracket M_1 \rrbracket_p \sqcup \llbracket M_2 \rrbracket_p) \llbracket M \rrbracket_p & \text{else if } p \in \text{pn}(M) \\ \llbracket M_1 \rrbracket_p \sqcup \llbracket M_2 \rrbracket_p & \text{otherwise} \end{cases} \\
\llbracket \text{com}_{q,r} \rrbracket_p &= \begin{cases} \text{send}_r & \text{if } p = q \\ \text{recv}_q & \text{if } p = r \\ \perp & \text{otherwise} \end{cases} \quad \llbracket f(p_1, \dots, p_n) \rrbracket_p = \begin{cases} f_i(p_1, \dots, p_{i-1}, p_{i+1}, \dots, p_n) & \text{if } p = p_i \\ \perp & \text{otherwise} \end{cases} \\
\llbracket \text{select}_{q,r} l M \rrbracket_p &= \begin{cases} \oplus_r l \llbracket M \rrbracket_q & \text{if } p = q \\ \&_q \{l : \llbracket M \rrbracket_r\} & \text{if } p = r \\ \llbracket M \rrbracket_p & \text{otherwise} \end{cases} \quad \llbracket c@r \rrbracket_p = \begin{cases} c & \text{if } p = r \\ \perp & \text{otherwise} \end{cases} \\
&\quad \llbracket x \rrbracket_p = \begin{cases} x & \text{if } p \in \text{pn}(\text{type}(x)) \\ \perp & \text{otherwise} \end{cases}
\end{aligned}$$

Operators:

$$\&_q \{l_i : P_i\}_{i \in \mathcal{I}} \sqcup \&_q \{l_j : P_j\}_{j \in \mathcal{J}} = \&_q \{l_k : P_k\}_{k \in \mathcal{I} \cup \mathcal{J}} \quad \text{if } \mathcal{I} \text{ and } \mathcal{J} \text{ are disjoint}$$

$$\perp \sqcup \perp = \perp$$

Types:

$$\llbracket T_1 \rightarrow_{\bar{p}} T_2 \rrbracket_p = \begin{cases} \llbracket T_1 \rrbracket_p \rightarrow \llbracket T_2 \rrbracket_p & \text{if } p \in \text{pn}(T_1) \cup \text{pn}(T_2) \cup \bar{p} \\ \perp & \text{otherwise} \end{cases} \quad \llbracket \alpha @ r \rrbracket_p = \begin{cases} \alpha & \text{if } p = r \\ \perp & \text{otherwise} \end{cases}$$

Fig. 9. Endpoint projection for λ^χ .

4.1 The Projection Function

Projection is formally defined in Figure 9 by the relation $\llbracket M \rrbracket_p = P$, where P is the projection of M for process p . The relation is defined by induction on the typing derivations of choreographies as in prior work [Cruz-Filipe et al. 2023; Hirsch and Garg 2022], though we omit the derivations in the figure for clarity. When the subscript is omitted, $\llbracket M \rrbracket$ is the parallel composition of the projections of the processes in the choreography, i.e., $\llbracket M \rrbracket = p_1[\llbracket M \rrbracket_{p_1}] \mid \dots \mid p_n[\llbracket M \rrbracket_{p_n}]$ where $\text{pn}(M) = \{p_1, \dots, p_n\}$.

Many of the cases in Figure 9 should be unsurprising. For example, $\text{com}_{p,q}$ is projected to send_q for p and recv_p for q . Notice also that, whenever a process p is not involved in M , the choreography is projected to the bottom value \perp . However, to understand the projection of an if-expression, readers unfamiliar with choreographic programming now need to be introduced to the concept of *projectability* [Cruz-Filipe and Montesi 2020; Montesi 2023]. Notice that Figure 9 contains the mapping:

$$\llbracket \text{if } M \text{ then } M_1 \text{ else } M_2 \rrbracket_p = \llbracket M_1 \rrbracket_p \sqcup \llbracket M_2 \rrbracket_p \quad \text{when } p \notin \text{pn}(M). \quad (6)$$

This is because, if the expression M is evaluated to a boolean at some q , there is no way for p to directly observe q 's value; we say that p needs *knowledge of choice*. To obtain knowledge of choice,

$$\begin{aligned}\mathcal{B} &::= \oplus_p l \mathcal{B} \mid \&_p \{l : \mathcal{B}\} \mid (\lambda x : S.\mathcal{B}) P \mid \bullet \\ \mathcal{O} &::= \mathcal{B}[\mathcal{O}] \mid \mathcal{E}[\mathcal{O}] \mid \bullet\end{aligned}$$

Fig. 10. Degenerate contexts for networks.

q must explicitly inform p about its decision via selections $\text{select}_{q,p} l_1$ in M_1 and $\text{select}_{q,p} l_2$ in M_2 , where l_1 and l_2 are distinct labels; these selections can be seen in Figure 4.

Choreographic languages ensure that knowledge of choice is propagated correctly via the *merge* operator (\sqcup) seen in Equation (6) and defined in Figure 9. The merge operator is a partial function, so the projection can be undefined if its criteria are not met. A choreography with no projection is said to be *unprojectable*. In the rest of this section, we will assume all choreographies are projectable unless otherwise stated. (Experts will notice our merge operator is more restrictive than in $\text{Chor}\lambda$; see Section 5 for details.)

4.2 Properties of Projection

We now establish some properties about projection and evaluation contexts, fulfilling the promises made in Section 3. All choreographies are assumed to be projectable and may contain free variables, unless otherwise stated.

Informally, the projection $\llbracket M \rrbracket_p$ is always “the same as M , but with all the parts not related to p edited out”. For example, values are projected to values and any M where p is not involved will be projected to \perp . The latter result is used to establish modularity: editing the behavior of q should not change the projection for p if $p \neq q$.

LEMMA 4.1. *If V is a choreographic value then $\llbracket V \rrbracket_p$ is a network-level value.*

LEMMA 4.2. *Let M be a choreography where $p \notin \text{pn}(M)$. Then $\llbracket M \rrbracket_p = \perp$.*

LEMMA 4.3 (MODULARITY). *A context C is a choreography with a unique hole \bullet in place of some subexpression. Let C be a context and M_1, M_2 choreographies such that $p \notin \text{pn}(M_1)$ and $p \notin \text{pn}(M_2)$. Then $\llbracket C[M_1] \rrbracket_p = \llbracket C[M_2] \rrbracket_p$.*

Below we establish novel properties of choreographic evaluation contexts. Lemma 4.4 is our “fundamental property of answering contexts”: if a process r is not involved in \mathcal{A} , then the context will disappear during projection. Likewise, Lemma 4.5 is the “fundamental property of choreographic frames”: if r is involved in the scrutinee of a choreographic frame \mathcal{F} , the projection will be a network-level frame, i.e., $\llbracket \mathcal{F}[M] \rrbracket_r = \mathcal{F}[\llbracket M \rrbracket_r]$ for some \mathcal{F} . Choreographic redexes Δ at r also project to network-level redexes δ at r , so long as Δ is not a (*commute*) or (*app*) redex, by Lemma 4.6. When the choreography is in normal form, these results imply that evaluation contexts containing redexes are projected to evaluation contexts containing redexes: $\llbracket \mathcal{E}_r[\Delta] \rrbracket_r = \mathcal{E}[\delta]$.

But before we can define these properties, there is a snag: we know $\llbracket \mathcal{E}_r[M] \rrbracket_p$ should be an evaluation context when $r = p$, but what about when $r \neq p$? For this we must introduce *degenerate contexts* \mathcal{B}, \mathcal{O} defined in Figure 10. These contexts will play an important role in Section 4.3, but for now we will simply assert that they emerge naturally in the lemmas below as a complement to evaluation contexts.

LEMMA 4.4. Let \mathcal{A} be an answering context.

- (1) If $r \notin \text{pn}(\mathcal{A})$ then $\llbracket \mathcal{A}[M] \rrbracket_r = \llbracket M \rrbracket_r$ for any M .
- (2) If $r \in \text{pn}(\mathcal{A})$ then there exists \mathcal{B} such that, for any M , $\llbracket \mathcal{A}[M] \rrbracket_r = \mathcal{B}[\llbracket M \rrbracket_r]$.

LEMMA 4.5. Let $M = \mathcal{F}[N]$ where $\Theta; \Gamma \vdash N : T$. Then there exists a network-level frame \mathcal{F} such that, for any N' where $\Theta; \Gamma \vdash N' : T$,

- (1) If $r \in \text{pn}(N')$ then $\mathcal{F}[\llbracket N' \rrbracket_r] = \llbracket \mathcal{F}[N'] \rrbracket_r$.
- (2) If $r \notin \text{pn}(N')$ then $\mathcal{F}[\llbracket N' \rrbracket_r] \xrightarrow{\tau} \llbracket \mathcal{F}[N'] \rrbracket_r$.

LEMMA 4.6. Let Δ be a redex at r that is not an (*app*) or (*commute*) redex. Then $\llbracket \Delta \rrbracket_r = \delta$ for some network-level redex δ .

LEMMA 4.7. Let $M = \mathcal{E}_r[N]$ where $\Theta; \Gamma \vdash N : T$. Then there exists a network-level evaluation context \mathcal{E} such that, for any N' where $\Theta; \Gamma \vdash N' : T$,

- (1) If $r \in \text{pn}(N')$ then $\mathcal{E}[\llbracket N' \rrbracket_r] = \llbracket \mathcal{E}_r[N'] \rrbracket_r$.
- (2) If $r \notin \text{pn}(N')$ then $\mathcal{E}[\llbracket N' \rrbracket_r] \xrightarrow{\tau} \llbracket \mathcal{E}_r[N'] \rrbracket_r$.

4.3 The Projection Theorem

How does a choreography M relate to its projection $\llbracket M \rrbracket$? One might hope they directly correspond, so $M \xrightarrow{\mu} M'$ implies $\llbracket M \rrbracket \xrightarrow{\mu} \llbracket M' \rrbracket$ and vice versa. As a pair of diagrams, where the vertical bar denotes endpoint projection:

$$\begin{array}{ccc} M & \xrightarrow{\mu} & M' \\ | & & | \\ \llbracket M \rrbracket & \xrightarrow{\mu} & \llbracket M' \rrbracket \end{array} \quad \begin{array}{ccc} M & \xrightarrow{\mu} & M' \\ | & & | \\ \llbracket M \rrbracket & \xrightarrow{\mu} & \llbracket M' \rrbracket \end{array}$$

This would mean the projection function $\llbracket - \rrbracket$ is a (*strong*) *bisimulation*. But unfortunately projection is not a bisimulation, even in first-order models. Let us consider several reasons why.

4.3.1 The Multistep Problem. Certain computations can be done in one step choreographically, but require multiple steps at the network level. Take substitution for example, with types omitted for legibility:

$$\begin{array}{ccc} \text{let } x = \text{comp}_{p,q} \text{ in } x & \xrightarrow{\tau} & \text{comp}_{p,q} \\ | & & | \\ p[(\lambda x. x) \text{ send}_q] \mid q[(\lambda x. x) \text{ recv}_p] & \xrightarrow{\tau} & p[\text{send}_q] \mid q[\text{recv}_p] \end{array}$$

Notice the projection in the bottom left corner is a pair of processes, with $\llbracket \text{let } x = \text{comp}_{p,q} \text{ in } x \rrbracket_p = (\lambda x. x) \text{ send}_q$ and $\llbracket \text{let } x = \text{comp}_{p,q} \text{ in } x \rrbracket_q = (\lambda x. x) \text{ recv}_p$ by the definition of projection in Figure 9. Whereas the choreography can reduce in one step using (*let*), the network must take two steps: a (*p-app*) step at p and one at q .

The textbook solution is to add runtime-only terms [Montesi 2023; Plyukhin et al. 2024]: in the example above, we could have the choreography M reduce to the “signposted” term $q.M$ to indicate that q has taken a step but p has not. Signposting would allow us to prove strong correspondence results between choreographies and networks, but it clutters the syntax and semantics. To make λ^X more accessible to newcomers, we avoid this route. We opt instead to exhibit a *weak bisimulation* between choreographies and their projections, so that for example $M \xrightarrow{\tau} M'$ implies $\llbracket M \rrbracket \xrightarrow{\tau} \llbracket M' \rrbracket$.

$$\begin{array}{ccccc}
\text{if true@p then } M_1 \text{ else } M_2 & \xrightarrow{\tau} & M_1 & \xrightarrow{\tau} & \text{true@q} \\
\downarrow & & & & \downarrow \\
p[[M]_p] \mid q[[M_1]_p \sqcup [M_2]_p] & \xrightarrow{-\tau} & p[[M_1]_p] \mid q[[M_1]_p \sqcup [M_2]_p] & \xrightarrow{-\tau} & p[\perp] \mid q[\text{true}]
\end{array}$$

4.3.3 *The Recursion Problem.* Choreographies in λ^X do not necessarily terminate. Consider the program (with types omitted):

$$M = \left(\begin{array}{l} \text{let } x = \text{diverge}(p) \text{ in} \\ \text{let } f = \text{comp}_{p,q} \text{ in} \\ f \end{array} \right),$$

$$\begin{aligned} \text{diverge}() &=_{\text{def}} \text{diverge}() \\ \llbracket M \rrbracket_{\mathbf{p}} &= (\lambda x. (\lambda f. f) \text{ send}_{\mathbf{q}}) \text{diverge}() \\ \llbracket M \rrbracket_{\mathbf{q}} &= (\lambda f. f) \text{recv}_{\mathbf{p}}. \end{aligned}$$
$$\begin{aligned} \llbracket M' \rrbracket_p &= (\lambda x. \text{send}_q) \text{diverge}() \\ \llbracket M' \rrbracket_q &= \text{recv}_p \end{aligned}$$

Chor λ solves the problem above by adding rewriting rules at the network level, allowing p to reduce the subterm $(\lambda f. f) \text{ send}_q \text{ out of order}$. These extra rules are safe to add because they have no side-effects and because the λ -calculus is confluent, but it is unsatisfying to have rules in the semantics of processes that are only used for proofs.

4.3.4 A Unified Approach. We present a novel approach that handles all three of the above problems in one fell swoop. The idea is to introduce a *prophecy* relation¹ on networks $N \rightsquigarrow N'$, where N' is obtained by pruning branches in N and performing τ -transitions—possibly out of order. Formally:

$$\begin{aligned}
 p[O[P]] \mid N &\rightsquigarrow p[O[P']] \mid N && \text{if } P \mapsto P' && (\text{compute}) \\
 p[O[\&q\{l_i : P_i\}_{i \in I}]] \mid N &\rightsquigarrow p[O[\&q\{l_i : P_i\}_{i \in J}]] \mid N && \text{if } J \subseteq I && (\text{prune}) \\
 p[O[\mathcal{F}[\mathcal{B}[P]]]] \mid N &\rightsquigarrow p[O[\mathcal{B}[\mathcal{F}[P]]]] \mid N && && (\text{commute}) \\
 p[O[\perp P]] \mid N &\rightsquigarrow p[O[(\lambda x : \perp. \perp) P]] \mid N && && (\text{bottom}_2)
 \end{aligned}$$

where O denotes a degenerate context, as defined in Figure 10. As we saw in Sections 4.3.1 to 4.3.3, some steps at the choreography level require prophecy steps at the network level:

LEMMA 4.8. *Let (\rightsquigarrow) be the reflexive transitive closure of (\rightsquigarrow) . If $M \xrightarrow{\tau} M'$ only by (commute) and (app) reductions, then $\llbracket M \rrbracket \rightsquigarrow \llbracket M' \rrbracket$.*

LEMMA 4.9. *Let p, r be roles such that $p \neq r$. Let $M = \mathcal{E}_r[N]$ where $\Theta; \Gamma \vdash N : T$. Then there exists a degenerate context O such that, for any N' where $\Theta; \Gamma \vdash N' : T$,*

- (1) *If $p \in \text{pn}(N')$ then $O[\llbracket N' \rrbracket_p] = \llbracket \mathcal{E}_r[N'] \rrbracket_p$.*
- (2) *If $p \notin \text{pn}(N')$ then $O[\llbracket N' \rrbracket_p] \rightsquigarrow \llbracket \mathcal{E}_r[N'] \rrbracket_p$.*

We will use the prophecy relation to capture the idea that a choreography “runs faster (and smarter!)” than its projection, but the projection can always take a finite number of prophecy steps to “catch up”. Formally: our goal will be to show that the relation $N \rightsquigarrow \llbracket M \rrbracket$ is a weak bisimulation between networks and choreographies.

We prove the result in two pieces. First, we prove the *Prophecy Theorem* (Theorem 4.10), which shows that prophecy steps commute with ordinary steps in the network. Second, we prove the *Projection Theorem* (Theorem 4.11), which shows that executions of choreographies and their choreographies correspond, mediated by the prophecy relation. Then we compose the two results to establish a weak bisimulation (Theorem 4.12). This proof technique is more modular than past work, which implicitly combines the Prophecy and Projection Theorems into one monolithic proof by induction [Cruz-Filipe et al. 2023; Montesi 2023].

We begin with the Prophecy Theorem. The completeness direction says that if a network N_1 can “catch up” to a choreography’s projection $\llbracket M \rrbracket$ using prophecy steps, then any step by $\llbracket M \rrbracket$ can be matched by N_1 after first performing some τ -transitions. Conversely, soundness means that any visible transition by N_1 can be matched by the projection, and invisible transitions by N_1 might not require the projection to take any steps at all.

THEOREM 4.10 (PROPHECY THEOREM). *Let M be a choreography, N a network, and $N \rightsquigarrow \llbracket M \rrbracket$.*

- (Completeness) *If $\llbracket M \rrbracket \xrightarrow{\mu} \tilde{N}'$ then there exists N' such that $N \xrightarrow{\tau} N' \rightsquigarrow \tilde{N}'$.*
- (Soundness) *If $N \xrightarrow{\mu} N'$ where $\mu \neq \tau$ then there exists \tilde{N}' such that $\llbracket M \rrbracket \xrightarrow{\mu} \tilde{N}'$ and $N' \rightsquigarrow \tilde{N}'$. If $N \xrightarrow{\tau} N'$ then either $N' \rightsquigarrow \llbracket M \rrbracket$ or there exists \tilde{N}' such that $\llbracket M \rrbracket \xrightarrow{\tau} \tilde{N}'$ and $N' \rightsquigarrow \tilde{N}'$.*

PROOF (SKETCH). Let $(\xrightarrow{\mu?})$ denote exactly one step if $\mu \neq \tau$ and at most one step if $\mu = \tau$. For the Completeness direction, it suffices to prove that if $N \rightsquigarrow \tilde{N} \xrightarrow{\mu} \tilde{N}'$ then there exists N' such that $N \xrightarrow{\tau} N' \rightsquigarrow \tilde{N}'$. For Soundness, it suffices to prove that if $N \rightsquigarrow \tilde{N} \rightsquigarrow \llbracket M \rrbracket$ and $N \xrightarrow{\mu} N'$ then

¹The name for this relation is inspired by the related idea of *prophecy variables* [Abadi and Lamport 1991].

there exists \tilde{N}' such that $\tilde{N} \xrightarrow{\mu^?} \tilde{N}'$ and $N' \rightsquigarrow \tilde{N}'$. As a pair of diagrams:

$$\begin{array}{ccc} \tilde{N} & \xrightarrow{\mu} & \tilde{N}' \\ \uparrow \wr & & \uparrow \wr \\ N & \xrightarrow{\tau} \rightsquigarrow \xrightarrow{\mu} & N' \end{array} \quad \begin{array}{ccc} \tilde{N} & \xrightarrow{\mu^?} & \tilde{N}' \\ \uparrow \wr & & \uparrow \wr \\ N & \xrightarrow{\mu} & N' \end{array}$$

The proofs proceed by straightforward pattern matching of critical pairs [Barendregt 1984]. \square

Next we establish the Projection Theorem. Unlike prior work, we mediate the correspondence between choreographies and their projection via the prophecy relation, and we use the properties of evaluation contexts established in Section 4.2. This extra structure leads to nice diagrammatic proofs, capturing our intuition for how the two models correspond.

The completeness direction tells us that compiled code exhibits all the behaviors visible in the choreography. The soundness direction (which is typically much harder to prove) tells us that compiled code *only* exhibits behaviors from the choreography.

THEOREM 4.11 (PROJECTION THEOREM). *Let M_1 be a closed choreography.*

- (Completeness) If $M_1 \xrightarrow{\mu} M_2$ is not an (app) or (commute) step, then $\llbracket M_1 \rrbracket \xrightarrow{\mu} \rightsquigarrow \llbracket M_2 \rrbracket$.
- (Soundness) If $\llbracket M_1 \rrbracket \xrightarrow{\mu} N_2$ then there exists M_2 such that $N_2 \rightsquigarrow \llbracket M_2 \rrbracket$ and $M_1 \xrightarrow{\tau} \xrightarrow{\mu} M_2$.

PROOF (COMPLETENESS). Proceed by case analysis on $M_1 \xrightarrow{\mu} M_2$.

- Assume $\mu = \text{comp}_{p,q} c$. Then $M_1 = \mathcal{E}_{p,q}[\text{comp}_{p,q} c]$. Since projection at p and q preserves evaluation contexts with redexes at p and q (Lemma 4.7), there exist $\mathcal{E}, \mathcal{E}', N', N_2$ such that:

$$\begin{array}{ccc} \mathcal{E}_{p,q}[\text{comp}_{p,q} c@p] & \xrightarrow{\text{comp}_{p,q} c} & \mathcal{E}_{p,q}[c@q] \\ \downarrow & & \downarrow \\ p[\mathcal{E}[\text{send}_q c]] \mid q[\mathcal{E}'[\text{recv}_p \perp]] \mid N' & \xrightarrow{\text{comp}_{p,q} c} N_2 \xrightarrow{\tau} \rightsquigarrow & p[\llbracket \mathcal{E}_{p,q}[c@q] \rrbracket_p] \mid q[\mathcal{E}'[c]] \mid N' \end{array}$$

Specifically, $N_2 = p[\mathcal{E}[\perp]] \mid q[\mathcal{E}'[c]] \mid N'$. The only complication is that $p \notin \text{pn}(c@q)$, so $\llbracket \mathcal{E}_{p,q}[c@q] \rrbracket_p \neq \mathcal{E}[\perp]$ in general; instead, Lemma 4.7 only guarantees process p can *catch up* to its projection via $\mathcal{E}[\perp] \xrightarrow{\tau} \llbracket \mathcal{E}_{p,q}[c@q] \rrbracket_p$. At every other process $r \neq \{p, q\}$, the projection is unchanged due to modularity (Lemma 4.3). Since $(\xrightarrow{\tau}) \subseteq (\rightsquigarrow)$, we have $\llbracket M_1 \rrbracket \xrightarrow{\mu} \rightsquigarrow \llbracket M_2 \rrbracket$.

- Assume $\mu = \text{select}_{p,q} l$. Similar to above.
- Assume $M_1 = \mathcal{E}_p[\text{let } x : T = V \text{ in } M]$ and that $M_2 = \mathcal{E}_p[M[x := V]]$. Crucially, one must show that all processes in $\text{pn}(V) \cup \text{pn}(M)$ can match the choreography's step—not just p. For illustrative purposes, take $q \in \text{pn}(V) \setminus \{p\}$. Then by Lemmas 4.1 and 4.7 there exist $N, \mathcal{E}, P, S, L, O, Q, L', S'$ such that:

$$\begin{array}{ccc} \mathcal{E}_p[\text{let } x : T = V \text{ in } M] & \xrightarrow{\tau} & \mathcal{E}_p[M[x := V]] \\ \downarrow & & \downarrow \\ p[\mathcal{E}[(\lambda x : S. P) L]] \mid q[O[(\lambda x : S'. Q) L']] \mid N & \xrightarrow{\tau} \rightsquigarrow & p[\mathcal{E}[P[x := L]]] \mid q[O[Q[x := L']]] \mid N \end{array}$$

To complete the commuting diagram, p may need to catch up to its projection like above, and q may need to use the prophecy relation (\rightsquigarrow). Below we enumerate the four possible cases, and see that $\llbracket M_1 \rrbracket_q \rightsquigarrow \llbracket M_2 \rrbracket_q$ for each q:

- For p, $\llbracket M_1 \rrbracket_p = \mathcal{E}[(\lambda x : \llbracket T \rrbracket_p. \llbracket M \rrbracket_p) \llbracket V \rrbracket_p]$. By a lemma in the Supplemental Material, $\llbracket M \rrbracket_q[x := \llbracket V \rrbracket_q] = \llbracket M[x := V] \rrbracket_q$. Hence, by Lemma 4.7, $\llbracket M_1 \rrbracket_p \xrightarrow{\tau} \llbracket M_2 \rrbracket_p$.

- For each $q \in \text{pn}(V) \setminus \{p\}$, there exists \mathcal{O} such that $\llbracket M_1 \rrbracket_q = \mathcal{O}[(\lambda x : \llbracket T \rrbracket_q. \llbracket M \rrbracket_q) \llbracket V \rrbracket_q]$ by Lemma 4.9. Then $\llbracket M_1 \rrbracket_q \rightsquigarrow \mathcal{O}[\llbracket M \rrbracket_q[x := \llbracket V \rrbracket_q]]$. Hence $\llbracket M_1 \rrbracket_q \rightsquigarrow \llbracket M_2 \rrbracket_q$.
- For each $q \in \text{pn}(M) \setminus \text{pn}(V)$, there exists \mathcal{O} such that $\llbracket M_1 \rrbracket_q = \mathcal{O}[\llbracket M \rrbracket_q]$ by Lemma 4.9. Hence $\llbracket M_1 \rrbracket_q = \llbracket M_2 \rrbracket_q$.
- For all remaining processes q , $\llbracket M_1 \rrbracket_q = \llbracket M_2 \rrbracket_q$ by Lemma 4.3.
- Assume $M_1 = \mathcal{E}_p[M]$ where $M = \text{if } c@p \text{ then } M_{\text{true}} \text{ else } M_{\text{false}}$. Without loss of generality, let $c = \text{true}$. Here we must show that every process $q \in \text{pn}(M_{\text{true}}) \cup \text{pn}(M_{\text{false}})$ can match the choreography's step. Using the same arguments as above, there exist $\mathcal{N}, \mathcal{E}, \mathcal{O}, P, I, \mathcal{J}$ where $\mathcal{J} \subseteq I$ and there exist l_i, Q_i for each $i \in I$ such that:

$$\begin{array}{ccc}
 \mathcal{E}_p[M] & \xrightarrow{\tau} & \mathcal{E}_p[M_{\text{true}}] \\
 | & & | \\
 p[\mathcal{E}[P]] \mid q[\mathcal{O}[\&_r\{l_i : Q_i\}_{i \in I}]] \mid \mathcal{N} & \xrightarrow{\tau} \rightsquigarrow & p[\mathcal{E}[P_{\text{true}}]] \mid q[\mathcal{O}[\&_r\{l_j : Q_j\}_{j \in \mathcal{J}}]] \mid \mathcal{N}
 \end{array}$$

Here q can match the choreography's step by pruning its branches with the prophecy relation.

- Assume $M_1 = \mathcal{E}_p[f(\bar{p})]$ and $M_2 = \mathcal{E}_p[M[\bar{p} := \bar{q}]]$, where $(f(\bar{q}) : S = M) \in \mathbb{D}$. By the same arguments as above, $\llbracket M_1 \rrbracket \xrightarrow{\tau} \rightsquigarrow \llbracket M_2 \rrbracket$.

□

PROOF (SOUNDNESS). Proof by case analysis on $\llbracket M_1 \rrbracket \xrightarrow{\mu} \mathcal{N}_2$. By Lemmas 3.2 and 4.8, we can assume M_1 is in normal form without loss of generality.

- Assume $\mu = \text{com}_{p,q}$. Then $\llbracket M_1 \rrbracket = p[\mathcal{E}[\text{send}_q c]] \mid q[\mathcal{E}'[\text{recv}_p \perp]] \mid \mathcal{N}'$. At this point in the proof, we do not know anything about the structure of M_1 —but we can use results established in Section 4.2 to discover that structure. By Lemmas 3.4, 4.1 and 4.4, there exist \mathcal{E}_p, Δ such that $M_1 = \mathcal{E}_p[\Delta]$. Moreover, since M_1 is in normal form, Δ cannot be a (*commute*) or (*app*) redex. By Lemmas 4.6 and 4.7, and the uniqueness of evaluation contexts in the network language, $\llbracket \mathcal{E}_p[\Delta] \rrbracket_p = \mathcal{E}[\llbracket \Delta \rrbracket_p]$ with $\llbracket \Delta \rrbracket_p = \text{send}_q c$. Hence Δ can only be $\text{com}_{p,q} c@p$. And by Lemma 3.5, \mathcal{E}_p must also be a choreographic evaluation context at q ; let us rename \mathcal{E}_p to $\mathcal{E}_{p,q}$. In summary, we have the following diagram:

$$\begin{array}{ccc}
 \mathcal{E}_{p,q}[\text{com}_{p,q} c@p] & \xrightarrow{\text{com}_{p,q} c} & \mathcal{E}_{p,q}[c@q] \\
 | & & | \\
 p[\mathcal{E}[\text{send}_q c]] \mid q[\mathcal{E}'[\text{recv}_p \perp]] \mid \mathcal{N}' & \xrightarrow{\text{com}_{p,q} c} \mathcal{N}_2 \xrightarrow{\tau} \rightsquigarrow & p[\llbracket \mathcal{E}_{p,q}[c@q] \rrbracket_p] \mid q[\mathcal{E}'[c]] \mid \mathcal{N}'
 \end{array}$$

In fact, this diagram is identical to the one we drew for the Completeness proof. Having established the correspondence between $\mathcal{E}_{p,q}$ and $\mathcal{E}, \mathcal{E}'$, the rest of the proof proceeds the same as it did then. The case for selections is similar.

- Assume $\mu = \tau$ and the step proceeds by reducing a (*p-app*) redex at p . Then we must have $\llbracket M_1 \rrbracket_p = \mathcal{E}[(\lambda x : S. P) L]$. By Lemmas 3.4, 4.6 and 4.7, there exist \mathcal{E}_p, T, M, V such that $M_1 = \mathcal{E}_p[\text{let } x : T = V \text{ in } M'_1]$. Letting $M_2 = \mathcal{E}_p[M'_1[x := V]]$, the rest of the proof proceeds the same as the Completeness proof.

- Assume $\mu = \tau$ and the step proceeds by reducing a $(p\text{-if})$ redex at p . Then we must have $\llbracket M_1 \rrbracket_p = \mathcal{E}[\text{if } c \text{ then } P_{\text{true}} \text{ else } P_{\text{false}}]$. Without loss of generality, let $c = \text{true}$. By the arguments as above, there exist $\mathcal{E}_p, M_{\text{true}}, M_{\text{false}}$ such that $M_1 = \mathcal{E}_p[\text{if true then } M_{\text{true}} \text{ else } M_{\text{false}}]$. This case now reduces to the same argument from the Completeness proof.
- Assume $\mu = \tau$ and the step proceeds by reducing a $(p\text{-def})$ redex at p . This case proceeds by the same arguments as above.

□

THEOREM 4.12. *The relation $\mathcal{N} \rightsquigarrow \llbracket M \rrbracket$ is a weak bisimulation.*

PROOF. (Completeness) Assume $\mathcal{N}_1 \rightsquigarrow \llbracket M_1 \rrbracket$ and $M_1 \xrightarrow{\mu} M_2$. If $M_1 \xrightarrow{\tau} M_2$ by (app) or $(commute)$, then $\llbracket M_1 \rrbracket \rightsquigarrow \llbracket M_2 \rrbracket$ by Lemma 4.8; letting $\mathcal{N}_1 = \mathcal{N}_2$, we trivially have $\mathcal{N}_1 \xrightarrow{\tau} \mathcal{N}_2$ and $\mathcal{N}_2 \rightsquigarrow \llbracket M_2 \rrbracket$. Otherwise, by the Projection Theorem there exists \mathcal{N}'_2 such that $\llbracket M_1 \rrbracket \xrightarrow{\mu} \mathcal{N}'_2 \rightsquigarrow \llbracket M_2 \rrbracket$. By the Prophecy Theorem, there exists \mathcal{N}_2 such that $\mathcal{N}_1 \xrightarrow{\tau} \xrightarrow{\mu} \mathcal{N}_2$ and $\mathcal{N}_2 \rightsquigarrow \mathcal{N}'_2$. Hence $\mathcal{N}_1 \xrightarrow{\tau} \xrightarrow{\mu} \xrightarrow{\tau} \mathcal{N}_2$ and $\mathcal{N}_2 \rightsquigarrow \llbracket M_2 \rrbracket$.

(Soundness) Assume $\mathcal{N}_1 \rightsquigarrow \llbracket M_1 \rrbracket$ and $\mathcal{N}_1 \xrightarrow{\mu} \mathcal{N}_2$. By the Prophecy Theorem, there are two cases. In the first case, $\mu = \tau$ and $\mathcal{N}_2 \rightsquigarrow \llbracket M_1 \rrbracket$; letting $M_1 = M_2$, we trivially have $M_1 \xrightarrow{\tau} M_2$ and $\mathcal{N}_2 \rightsquigarrow \llbracket M_2 \rrbracket$. In the second case, there exists \mathcal{N}'_2 such that $\mathcal{N}_2 \rightsquigarrow \mathcal{N}'_2$ and $\llbracket M_1 \rrbracket \xrightarrow{\mu} \mathcal{N}'_2$. By the Projection Theorem, there exists M_2 such that $M_1 \xrightarrow{\tau} \xrightarrow{\mu} M_2$ and $\mathcal{N}'_2 \rightsquigarrow \llbracket M_2 \rrbracket$. Hence $M_1 \xrightarrow{\tau} \xrightarrow{\mu} \xrightarrow{\tau} M_2$ and $\mathcal{N}_2 \rightsquigarrow \llbracket M_2 \rrbracket$. □

With the prophecy relation, we eliminated the need for restructuring rules at the network level and made it convenient to prove the Projection Theorem by case analysis on evaluation contexts. Experts in choreographic programming will know that proving the Projection Theorem typically requires a large and tedious induction proof with many similar cases; our proof technique reduces the burden to just a few diagrams with no need for induction. We conclude with a victory lap: proving deadlock-freedom by design.

THEOREM 4.13 (DEADLOCK-FREEDOM). *Let M_0 be a choreography with $\text{pn}(M_0) = \{p_1, \dots, p_n\}$. If $\llbracket M_0 \rrbracket$ evaluates to a network \mathcal{N} that cannot be evaluated any further, then $\mathcal{N} = p_1[L_1] \mid \dots \mid p_n[L_n]$ where L_1, \dots, L_n are all values.*

PROOF. By Theorem 4.12, M_0 evaluates to some M such that $\mathcal{N} \rightsquigarrow \llbracket M \rrbracket$. Let \tilde{M} be the normal form of M , so that $M \xrightarrow{\tau} \tilde{M}$; by Lemma 4.8, $\llbracket M \rrbracket \rightsquigarrow \llbracket \tilde{M} \rrbracket$ and therefore $\mathcal{N} \rightsquigarrow \llbracket \tilde{M} \rrbracket$. Notice \tilde{M} cannot be evaluated any further: by the Projection Theorem, another step $\tilde{M} \xrightarrow{\mu} M'$ would imply $\llbracket \tilde{M} \rrbracket \xrightarrow{\mu} \rightsquigarrow \llbracket M' \rrbracket$, and so the Prophecy Theorem would imply $\mathcal{N} \xrightarrow{\tau} \xrightarrow{\mu} \mathcal{N}'$ for some \mathcal{N}' . Since \tilde{M} cannot be evaluated further, Theorem 3.6 implies \tilde{M} is a choreographic value V . Hence, by Lemma 4.1, $\llbracket \tilde{M} \rrbracket_{p_i}$ is a value L_i for each $p_i \in \text{pn}(M)$. Finally, by inspecting the definition of (\rightsquigarrow) , $\mathcal{N} \rightsquigarrow \llbracket \tilde{M} \rrbracket$ is only possible if $\mathcal{N} = \llbracket \tilde{M} \rrbracket$. □

5 Related Work

Higher-Order Choreographies. Higher-order choreographic programming was introduced by the Choral programming language [Giallorenzo et al. 2020, 2024]. Choral is also the language that introduced the idea of modeling choreographic data structures and communication by extending mainstream data types with locations and then having functions that input and output data at different locations. The theoretical foundations of this idea have been investigated in Chorλ [Cruz-Filipe et al. 2022; Cruz-Filipe et al. 2023] and Pirouette [Hirsch and Garg 2022]. However, as we

explained in Section 1, these models sacrifice either adequacy or elegance; λ^x addresses both. Note that, although λ^x is based on $\text{Chor}\lambda$, in principle we could develop the same results with *Pirouette* by changing its evaluation strategy.

Multiparty session types support “nested protocols”, which can be seen as a form of higher-order composition for simple choreographies without computation [Demangeon and Honda 2012]. Differently from most higher-order choreographic programming languages and our approach, where code is fully concurrent via out-of-order execution, nested multiparty session types require fixing a role that acts as an orchestrator to direct when a sub-choreography is entered.

Other Models. Choreographies in our model have two slightly unusual features: explicit let-bindings and a special (*commute*) rule. Both features have a long history in past work.

Many readers will be familiar with Moggi’s computational λ -calculus, where let-bindings express the sequencing of effects [Moggi 1991]. But explicit let-bindings also arise when embedding λ -calculi into proof nets, where $(\lambda x.M) N \mapsto \text{let } x = N \text{ in } M$ and $\text{let } x = V \text{ in } M \mapsto M[x := V]$ correspond to multiplicative and exponential cut-elimination, respectively [Accattoli 2015]. Ariola et al. [1995] [1989] also used let-bindings to model sharing in non-strict calculi, similarly to our model.

Most of the models above also require restructuring rules, like those in Figure 8. For instance, our (*let-let*) rule corresponds to monad associativity [Moggi 1991]; our (*let-app*) rule is needed for sharing in call-by-need [Ariola et al. 1995]; our (*let-app*) and (*app-let*) rules are needed to internally characterize solvability and achieve other good properties in call-by-value [Carraro and Guerrieri 2014; Herbelin and Zimmermann 2009]; and similar rules arise when embedding call-by-name terms into proof nets [Régnier 1994]. Maurer et al. [2017] observed that all the rules above can be neatly summarized by a single axiom, which pushes frames inside answering contexts. In choreographic models like λ^x and $\text{Chor}\lambda$, the (*commute*) rule ensures soundness, i.e., choreographies can exhibit all the same behaviors as their projections.

Distributed Data Types. The full version of $\text{Chor}\lambda$ includes constructors for distributed products and sums—we omitted these constructors in pursuit of a simple model. As usual, one can partially recover these constructors with lambda encodings. Consider the standard encodings for products and sums, augmented with extra process annotations:

$$\begin{aligned}
 \text{Pair} &\equiv \lambda x_1 : T_1. \lambda x_2 : T_2. \lambda p : T_1 \rightarrow T_2 \rightarrow_{\bar{p}} T. p \ x_1 \ x_2 \\
 \text{fst} &\equiv \lambda x_1 : T_1. \lambda x_2 : T_2. x_1 \\
 \text{snd} &\equiv \lambda x_1 : T_1. \lambda x_2 : T_2. x_2 \\
 \text{Inl} &\equiv \lambda x : T_1. \lambda l : T_1 \rightarrow_{\bar{p}} T. \lambda r : T_2 \rightarrow_{\bar{q}} T. l \ x \\
 \text{Inr} &\equiv \lambda x : T_2. \lambda l : T_1 \rightarrow_{\bar{p}} T. \lambda r : T_2 \rightarrow_{\bar{q}} T. r \ x \\
 \text{case} &\equiv \lambda s : (T_1 \rightarrow_{\bar{p}} T) \rightarrow (T_2 \rightarrow_{\bar{q}} T) \rightarrow_{\bar{p} \cup \bar{q}} T. \lambda l : T_1 \rightarrow_{\bar{p}} T. \lambda r : T_2 \rightarrow_{\bar{q}} T. s \ l \ r
 \end{aligned}$$

The principal limitation of this encoding is that the term constructing the datatype must dictate in advance what the type T of the handler will be. In λ^x , this also means fixing a *location* for the type T and a set of *proxy processes* \bar{p} involved in computing T . We can loosen these restrictions by adding explicit support for ADTs, as in $\text{Chor}\lambda$, or by adding process polymorphism, as in $\text{PolyChor}\lambda$ [Graversen et al. 2024].

Selections. HasChor [Shen et al. 2023] is a library for choreographic programming in Haskell. In HasChor, if-expressions and selections are merged into a single construct. When one writes *if* $f(p)$ *then* M_1 *else* M_2 , process p implicitly sends a selection to every other process in the choreography—including processes not involved in M_1 or M_2 . Choreographic conclaves [Bates et al. 2025] improve on this mechanism, ensuring only processes involved in the branches will receive a

selection. These approaches are useful in library-level choreographic programming, where it is difficult to statically check knowledge of choice.

However, explicit selections are more common in choreographic programming because they offer more control. For instance, in λ^χ we may write:

$$\begin{array}{l} \text{if } f(p) \\ \text{then } \text{select}_{p,q} L_1 M_1 \\ \text{else } \left(\begin{array}{l} \text{if } g(p) \\ \text{then } \text{select}_{p,q} L_2 M_2 \\ \text{else } \text{select}_{p,q} L_3 M_3 \end{array} \right) \end{array}$$

which projects to the following process at q :

$$\&_q \{L_1 : \llbracket M_1 \rrbracket_q, L_2 : \llbracket M_2 \rrbracket_q, L_3 : \llbracket M_3 \rrbracket_q\}$$

This term only requires one selection from p to q , whereas in HasChor it would require two selections—one for each if-expression. There has been significant work on making explicit selections even more powerful: [Lugovic and Montesi 2024] combines selections and communications into a single “type-driven communication”; [Cruz-Filipe et al. 2023b] coalesces all the selections of a loop into one message; and [Cruz-Filipe and Montesi 2023] develops an algorithm for inferring selections in first-order choreographies.

Evaluation Under Conditionals. One rule of Chor λ we decided not to capture is *InCase*, which permits execution underneath an if-expression. This rule is also present in Montesi’s textbook on choreographic programming, under the name *DelayCond* [Montesi 2023]. The inclusion of *InCase* is typically justified by a choreography like the following:

$$\text{if } c@p \text{ then } \text{com}_{p,q} 0@p \text{ else } \text{com}_{p,q} 1@p,$$

This choreography is not projectable in λ^χ , but it is projectable in Chor λ and can be implemented by the network $p[\text{if } c \text{ then } \text{send}_q 0 \text{ else } \text{send}_q 1] \mid q[\text{recv}_p \perp]$. However, in simple examples like these we can make the choreography projectable in λ^χ by “pulling out” the communication:

$$\text{com}_{p,q} (\text{if } c@p \text{ then } 0@p \text{ else } 1@p)$$

The future of the *InCase* rule is unclear. Library-level choreographic programming languages omit the rule because it would be difficult to encode in the host language [Bates et al. 2025; Shen et al. 2023]; recent work in fully out-of-order choreographies² is incompatible with the rule as it is currently understood [Plyukhin et al. 2024]; and recent extensions like multiply-located values [Bates et al. 2025] could further increase the expressiveness of λ^χ without requiring *InCase*.

6 Conclusion

We have presented λ^χ , an elegant model for higher-order choreographic programming based on the λ -calculus. In particular, the model shows deep connections between choreographies and non-strict λ -calculi, culminating in a new evaluation strategy we call *semilenient* evaluation. We hope this work has made higher-order choreographies more accessible, particularly for experts from other fields. We conclude by discussing applications of λ^χ and opportunities for future work.

²In our terminology, “fully” out-of-order execution means giving choreographies a *lenient* semantics instead of *semilenient*.

Implementing EPP and source-level reasoning. To implement a higher-order choreographic language, compiler authors need to choose a projection function. If they use projection *à la* Pirouette, the language will have an ordinary call-by-value semantics—but only if processes synchronize globally when they enter choreographic procedures. The compiler author could choose to omit these synchronizations, but then the semantics is unknown: users will have no way to predict their programs’ behavior except by inspecting the compiled endpoint code. To generate efficient code and retain a connection to some formal model, the compiler author could instead use projection *à la* Chor λ . But this is not much better, because users need to understand all Chor λ ’s unusual rules and edge cases before they can understand their programs.

Our model simplifies Chor λ , revealing that its evaluation strategy is straightforward after all: it is just *semilenient*, instead of the usual call-by-value. This means compiler authors can omit needless global synchronizations guilt-free, and compiler users can reason about their programs using our simplified model. Compiler authors could also use our model to develop choreography-level optimizations, like eliminating unnecessary communications, without changing the behavior of the user’s program.

Compiler testing. How do we know the code generated by a choreographic compiler is correct? Our Projection Theorem tells us what programmers should be able to expect: if M is a choreography that reduces to value V and the compiler is correct, then the code generated by the compiler should evaluate to $\llbracket V \rrbracket_p$ at each role p . Moreover, the order of communications we observe in the compiled code should correspond to some execution of M in the choreography semantics. This is a standard result that we can do with any choreography model.

But the principled design of our model also suggests *syntax-directed* ways to test a choreographic compiler. For example, we can test that Lemma 4.4 holds: answering contexts at p should disappear in the projection at p . We can also test Lemma 4.5: frames are projected into frames. If these properties hold, it suggests the compiler does not introduce any unintentional synchronization that would hurt performance.

A principled foundation for new languages. Choreographic programming is a very active field of study. Researchers and hobbyists alike are developing choreographic languages with novel features; these features are implemented by starting from an existing model, adding syntax, updating the type system, and implementing projection. But how can we gain confidence that the resulting language retains important properties like deadlock-freedom and type-safety? With prior work, we could only gain this confidence by first formalizing a new semantics *ex nihilo* and proving the projection theorem. Doing this is tedious and error-prone for researchers, and out of the question for working programmers.

Our model gives compiler authors *design principles* for new choreographic languages. We have argued already why we expect the semantics to be semilenient. We have also shown that semilenient semantics emerges naturally from the definition of answering contexts and choreographic frames. These two constructs have formal properties that are easy to check, c.f. Lemmas 4.4 and 4.5. Thus, language designers can already start to have confidence in new features by identifying appropriate generalizations of answering contexts and choreographic frames, and then checking that their projection algorithm satisfies those properties. Eventually, given the similarity of existing choreographic languages to one another, researchers may develop tools that automate the tedious proofs of properties like deadlock-freedom and the projection theorem.

Future work. We argued indirectly why λ^X is a good foundation, but the proof is in the pudding. How well does our approach play with orthogonal extensions, like process polymorphism [Graversen

et al. 2024], census polymorphism and multiply-located values [Bates et al. 2025], and fully out-of-order execution [Plyukhin et al. 2024]? Is the machinery we introduced sufficient for more complex models, and can it guide researchers toward the “right” abstractions?

Reducing proof burden is another important topic for future work. Although our approach makes proofs more “modular”, one can still easily make mistakes. Perhaps the building blocks we introduced here could be factored into reusable proofs or tactics in proof assistants like Rocq or Lean.

Acknowledgments

Partially supported by Villum Fonden (grant no. 29518). Co-funded by the European Union (ERC, CHORDS, 101124225). Views and opinions expressed are however those of the authors only and do not necessarily reflect those of the European Union or the European Research Council. Neither the European Union nor the granting authority can be held responsible for them.

References

2025. Chorex: Choreographic Programming in Elixir. <https://github.com/utahplt/chorex>.
2025. Klor: Choreographies in Clojure. <https://github.com/lovrodsu/klor>.
- Martín Abadi and Leslie Lamport. 1991. The Existence of Refinement Mappings. *Theor. Comput. Sci.* 82, 2 (1991), 253–284. doi:10.1016/0304-3975(91)90224-P
- Cosku Acay, Rolph Recto, Joshua Gancher, Andrew C. Myers, and Elaine Shi. 2021. Viaduct: an extensible, optimizing compiler for secure distributed programs. In *PLDI '21: 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation, Virtual Event, Canada, June 20-25, 2021*, Stephen N. Freund and Eran Yahav (Eds.). ACM, 740–755. doi:10.1145/3453483.3454074
- Beniamino Accattoli. 2015. Proof Nets and the Call-by-Value Lambda-Calculus. *Theor. Comput. Sci.* 606 (2015), 2–24. doi:10.1016/j.TCS.2015.08.006
- Zena Ariola and Arvind. 1989. P-TAC: A Parallel Intermediate Language. In *Proceedings of the Fourth International Conference on Functional Programming Languages and Computer Architecture - FPCA '89*. ACM Press, Imperial College, London, United Kingdom, 230–242. doi:10.1145/99370.99388
- Zena M. Ariola, John Maraist, Martin Odersky, Matthias Felleisen, and Philip Wadler. 1995. A Call-by-Need Lambda Calculus. In *Proceedings of the 22nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages - POPL '95*. ACM Press, San Francisco, California, United States, 233–246. doi:10.1145/199448.199507
- Arvind, Jan-Willem Maessen, Rishiyur S. Nikhil, and Joseph Stoy. 1996. A Lambda Calculus with Letrecs and Barriers. In *Foundations of Software Technology and Theoretical Computer Science*, Gerhard Goos, Juris Hartmanis, Jan Leeuwen, V. Chandru, and V. Vinay (Eds.). Vol. 1180. Springer Berlin Heidelberg, Berlin, Heidelberg, 19–36. doi:10.1007/3-540-62034-6_34
- Arvind, Rishiyur S. Nikhil, and Keshav Pingali. 1986. I-Structures: Data Structures for Parallel Computing. In *Graph Reduction, Proceedings of a Workshop, Santa Fé, New Mexico, USA, September 29 - October 1, 1986 (Lecture Notes in Computer Science, Vol. 279)*, Joseph H. Fasel and Robert M. Keller (Eds.). Springer, 336–369. doi:10.1007/3-540-18420-1_65
- Lennart Augustsson, Joachim Breitner, Koen Claessen, Ranjit Jhala, Simon Peyton Jones, Olin Shivers, Guy L. Steele Jr., and Tim Sweeney. 2023. The Verse Calculus: A Core Calculus for Deterministic Functional Logic Programming. *Proc. ACM Program. Lang.* 7, ICFP (2023), 417–447. doi:10.1145/3607845
- Hendrik Pieter Barendregt. 1984. *The Lambda Calculus: Its Syntax and Semantics* (rev. ed ed.). Number v. 103 in Studies in Logic and the Foundations of Mathematics. North-Holland Sole distributors for the U.S.A. and Canada, Elsevier Science Pub. Co, Amsterdam New York New York, N.Y.
- Mako Bates, Shun Kashiwa, Syed Jafri, Gan Shen, Lindsey Kuper, and Joseph P. Near. 2025. Efficient, Portable, Census-Polymorphic Choreographic Programming. In *PLDI '25: 46th ACM SIGPLAN International Conference on Programming Language Design and Implementation*. ACM, Seoul, South Korea. doi:10.1145/3729296
- Marco Carbone and Fabrizio Montesi. 2013. Deadlock-freedom-by-design: multiparty asynchronous global programming. In *The 40th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL '13, Rome, Italy - January 23 - 25, 2013*, Roberto Giacobazzi and Radhia Cousot (Eds.). ACM, 263–274. doi:10.1145/2429069.2429101
- Marco Carbone, Fabrizio Montesi, and Carsten Schürmann. 2018. Choreographies, logically. *Distributed Comput.* 31, 1 (2018), 51–67. doi:10.1007/S00446-017-0295-1
- Alberto Carraro and Giulio Guerrieri. 2014. A Semantical and Operational Account of Call-by-Value Solvability. In *Foundations of Software Science and Computation Structures - 17th International Conference, FOSSACS 2014, Held as Part of*

- the European Joint Conferences on Theory and Practice of Software, ETAPS 2014, Grenoble, France, April 5-13, 2014, *Proceedings (Lecture Notes in Computer Science, Vol. 8412)*, Anca Muscholl (Ed.). Springer, 103–118. doi:10.1007/978-3-642-54830-7_7
- Luís Cruz-Filipe, Eva Graversen, Lovro Lugovic, Fabrizio Montesi, and Marco Peressotti. 2022. Functional Choreographic Programming. In *Theoretical Aspects of Computing - ICTAC 2022 - 19th International Colloquium, Tbilisi, Georgia, September 27-29, 2022, Proceedings (Lecture Notes in Computer Science, Vol. 13572)*, Helmut Seidl, Zhiming Liu, and Corina S. Pasareanu (Eds.). Springer, 212–237. doi:10.1007/978-3-031-17715-6_15
- Luís Cruz-Filipe, Eva Graversen, Lovro Lugović, Fabrizio Montesi, and Marco Peressotti. 2023. Modular Compilation for Higher-Order Functional Choreographies. In *37th European Conference on Object-Oriented Programming, ECOOP 2023, July 17-21, 2023, Seattle, Washington, United States (LIPIcs, Vol. 263)*, Karim Ali and Guido Salvaneschi (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 7:1–7:37. doi:10.4230/LIPICS.ECOOP.2023.7
- Luís Cruz-Filipe, Eva Graversen, Fabrizio Montesi, and Marco Peressotti. 2023a. Reasoning About Choreographic Programs. In *Coordination Models and Languages - 25th IFIP WG 6.1 International Conference, COORDINATION 2023, Held as Part of the 18th International Federated Conference on Distributed Computing Techniques, DisCoTec 2023, Lisbon, Portugal, June 19-23, 2023, Proceedings (Lecture Notes in Computer Science, Vol. 13908)*, Sung-Shik Jongmans and Antónia Lopes (Eds.). Springer, 144–162. doi:10.1007/978-3-031-35361-1_8
- Luís Cruz-Filipe and Fabrizio Montesi. 2020. A Core Model for Choreographic Programming. *Theor. Comput. Sci.* 802 (2020), 38–66. doi:10.1016/j.tcs.2019.07.005
- Luís Cruz-Filipe and Fabrizio Montesi. 2023. Now It Compiles! Certified Automatic Repair of Uncompilable Protocols. In *14th International Conference on Interactive Theorem Proving, ITP 2023, July 31 to August 4, 2023, Białystok, Poland (LIPIcs, Vol. 268)*, Adam Naumowicz and René Thiemann (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 11:1–11:19. doi:10.4230/LIPICS.ITP.2023.11
- Luís Cruz-Filipe, Fabrizio Montesi, and Robert R. Rasmussen. 2023b. Keep me out of the loop: a more flexible choreographic projection. In *LPAR 2023: Proceedings of 24th International Conference on Logic for Programming, Artificial Intelligence and Reasoning, Manizales, Colombia, 4-9th June 2023 (EPIC Series in Computing, Vol. 94)*, Ruzica Piskac and Andrei Voronkov (Eds.). EasyChair, 144–163. doi:10.29007/WBW3
- Romain Demangeon and Kohei Honda. 2012. Nested Protocols in Session Types. In *CONCUR 2012 - Concurrency Theory - 23rd International Conference, CONCUR 2012, Newcastle upon Tyne, UK, September 4-7, 2012. Proceedings (Lecture Notes in Computer Science, Vol. 7454)*, Maciej Koutny and Irek Ulidowski (Eds.). Springer, 272–286. doi:10.1007/978-3-642-32940-1_20
- Cormac Flanagan, Amr Sabry, Bruce F. Duba, and Matthias Felleisen. 1993. The Essence of Compiling with Continuations. *ACM SIGPLAN Notices* 28, 6 (June 1993), 237–247. doi:10.1145/173262.155113
- Saverio Giallorenzo, Fabrizio Montesi, and Marco Peressotti. 2020. Choreographies as Objects. *CoRR abs/2005.09520* (2020). arXiv:2005.09520 <https://arxiv.org/abs/2005.09520>
- Saverio Giallorenzo, Fabrizio Montesi, and Marco Peressotti. 2024. Choral: Object-oriented Choreographic Programming. *ACM Trans. Program. Lang. Syst.* 46, 1 (2024), 1:1–1:59. doi:10.1145/3632398
- Eva Graversen, Andrew K. Hirsch, and Fabrizio Montesi. 2024. Alice or Bob?: Process Polymorphism in Choreographies. *Journal of Functional Programming* 34 (Jan. 2024). doi:10.1017/S0956796823000114
- Hugo Herbelin and Stéphane Zimmermann. 2009. An Operational Account of Call-by-Value Minimal and Classical Lambda-Calculus in "Natural Deduction" Form. In *Typed Lambda Calculi and Applications, 9th International Conference, TLCA 2009, Brasília, Brazil, July 1-3, 2009. Proceedings (Lecture Notes in Computer Science, Vol. 5608)*, Pierre-Louis Curien (Ed.). Springer, 142–156. doi:10.1007/978-3-642-02273-9_12
- Andrew K. Hirsch and Deepak Garg. 2022. Pirouette: higher-order typed functional choreographies. *Proc. ACM Program. Lang.* 6, POPL (2022), 1–27. doi:10.1145/3498684
- Kohei Honda, Nobuko Yoshida, and Marco Carbone. 2016. Multiparty Asynchronous Session Types. *J. ACM* 63, 1 (2016), 9:1–9:67. doi:10.1145/2827695
- Gérard Huet. 1997. The Zipper. *Journal of Functional Programming* 7, 5 (1997), 549–554. doi:10.1017/S0956796897002864
- Shadaj Laddad, Alvin Cheung, and Joseph M. Hellerstein. 2024. Suki: Choreographed Distributed Dataflow in Rust. arXiv:2406.14733 [cs]
- Lovro Lugovic and Fabrizio Montesi. 2024. Real-World Choreographic Programming: Full-Duplex Asynchrony and Interoperability. *Art Sci. Eng. Program.* 8, 2 (2024). doi:10.22152/PROGRAMMING-JOURNAL.ORG/2024/8/8
- Luke Maurer, Paul Downen, Zena M. Ariola, and Simon L. Peyton Jones. 2017. Compiling without Continuations. In *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI 2017, Barcelona, Spain, June 18-23, 2017, Albert Cohen and Martin T. Vechev (Eds.)*. ACM, 482–494. doi:10.1145/3062341.3062380
- Eugenio Moggi. 1991. Notions of Computation and Monads. *Information and Computation* 93, 1 (July 1991), 55–92. doi:10.1016/0890-5401(91)90052-4
- Fabrizio Montesi. 2013. *Choreographic Programming*. Ph.D. Thesis. IT University of Copenhagen. <https://www.fabriziomontesi.com/files/choreographic-programming.pdf>.

- Fabrizio Montesi. 2023. *Introduction to Choreographies*. Cambridge University Press, Cambridge.
- Roger M. Needham and Michael D. Schroeder. 1978. Using Encryption for Authentication in Large Networks of Computers. *Commun. ACM* 21, 12 (1978), 993–999. doi:[10.1145/359657.359659](https://doi.org/10.1145/359657.359659)
- Object Management Group. 2017. Unified Modeling Language, Version 2.5.1. <https://www.omg.org/spec/UML/2.5.1/PDF>.
- Gordon D. Plotkin. 1975. Call-by-Name, Call-by-Value and the lambda-Calculus. *Theor. Comput. Sci.* 1, 2 (1975), 125–159. doi:[10.1016/0304-3975\(75\)90017-1](https://doi.org/10.1016/0304-3975(75)90017-1)
- Dan Plyukhin, Marco Peressotti, and Fabrizio Montesi. 2024. Ozone: Fully Out-of-Order Choreographies. In *38th European Conference on Object-Oriented Programming, ECOOP 2024, September 16–20, 2024, Vienna, Austria (LIPIcs, Vol. 313)*, Jonathan Aldrich and Guido Salvaneschi (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 31:1–31:28. doi:[10.4230/LIPICS.ECOOP.2024.31](https://doi.org/10.4230/LIPICS.ECOOP.2024.31)
- Laurent Régnier. 1994. Une Équivalence Sur Les Lambda-Terms. *Theoretical Computer Science* 126, 2 (April 1994), 281–292. doi:[10.1016/0304-3975\(94\)90012-4](https://doi.org/10.1016/0304-3975(94)90012-4)
- Gan Shen, Shun Kashiwa, and Lindsey Kuper. 2023. HasChor: Functional Choreographic Programming for All (Functional Pearl). *Proc. ACM Program. Lang.* 7, ICFP (2023), 541–565. doi:[10.1145/3607849](https://doi.org/10.1145/3607849)
- Guy Tremblay. 2000. Lenient Evaluation Is Neither Strict nor Lazy. *Comput. Lang.* 26, 1 (2000), 43–66. doi:[10.1016/S0096-0551\(01\)00006-6](https://doi.org/10.1016/S0096-0551(01)00006-6)
- Petra van den Bos and Sung-Shik Jongmans. 2023. VeyMont: Parallelising Verified Programs Instead of Verifying Parallel Programs. In *Formal Methods - 25th International Symposium, FM 2023, Lübeck, Germany, March 6–10, 2023, Proceedings (Lecture Notes in Computer Science, Vol. 14000)*, Marsha Chechik, Joost-Pieter Katoen, and Martin Leucker (Eds.). Springer, 321–339. doi:[10.1007/978-3-031-27481-7_19](https://doi.org/10.1007/978-3-031-27481-7_19)

Received 2025-02-27; accepted 2025-06-27